

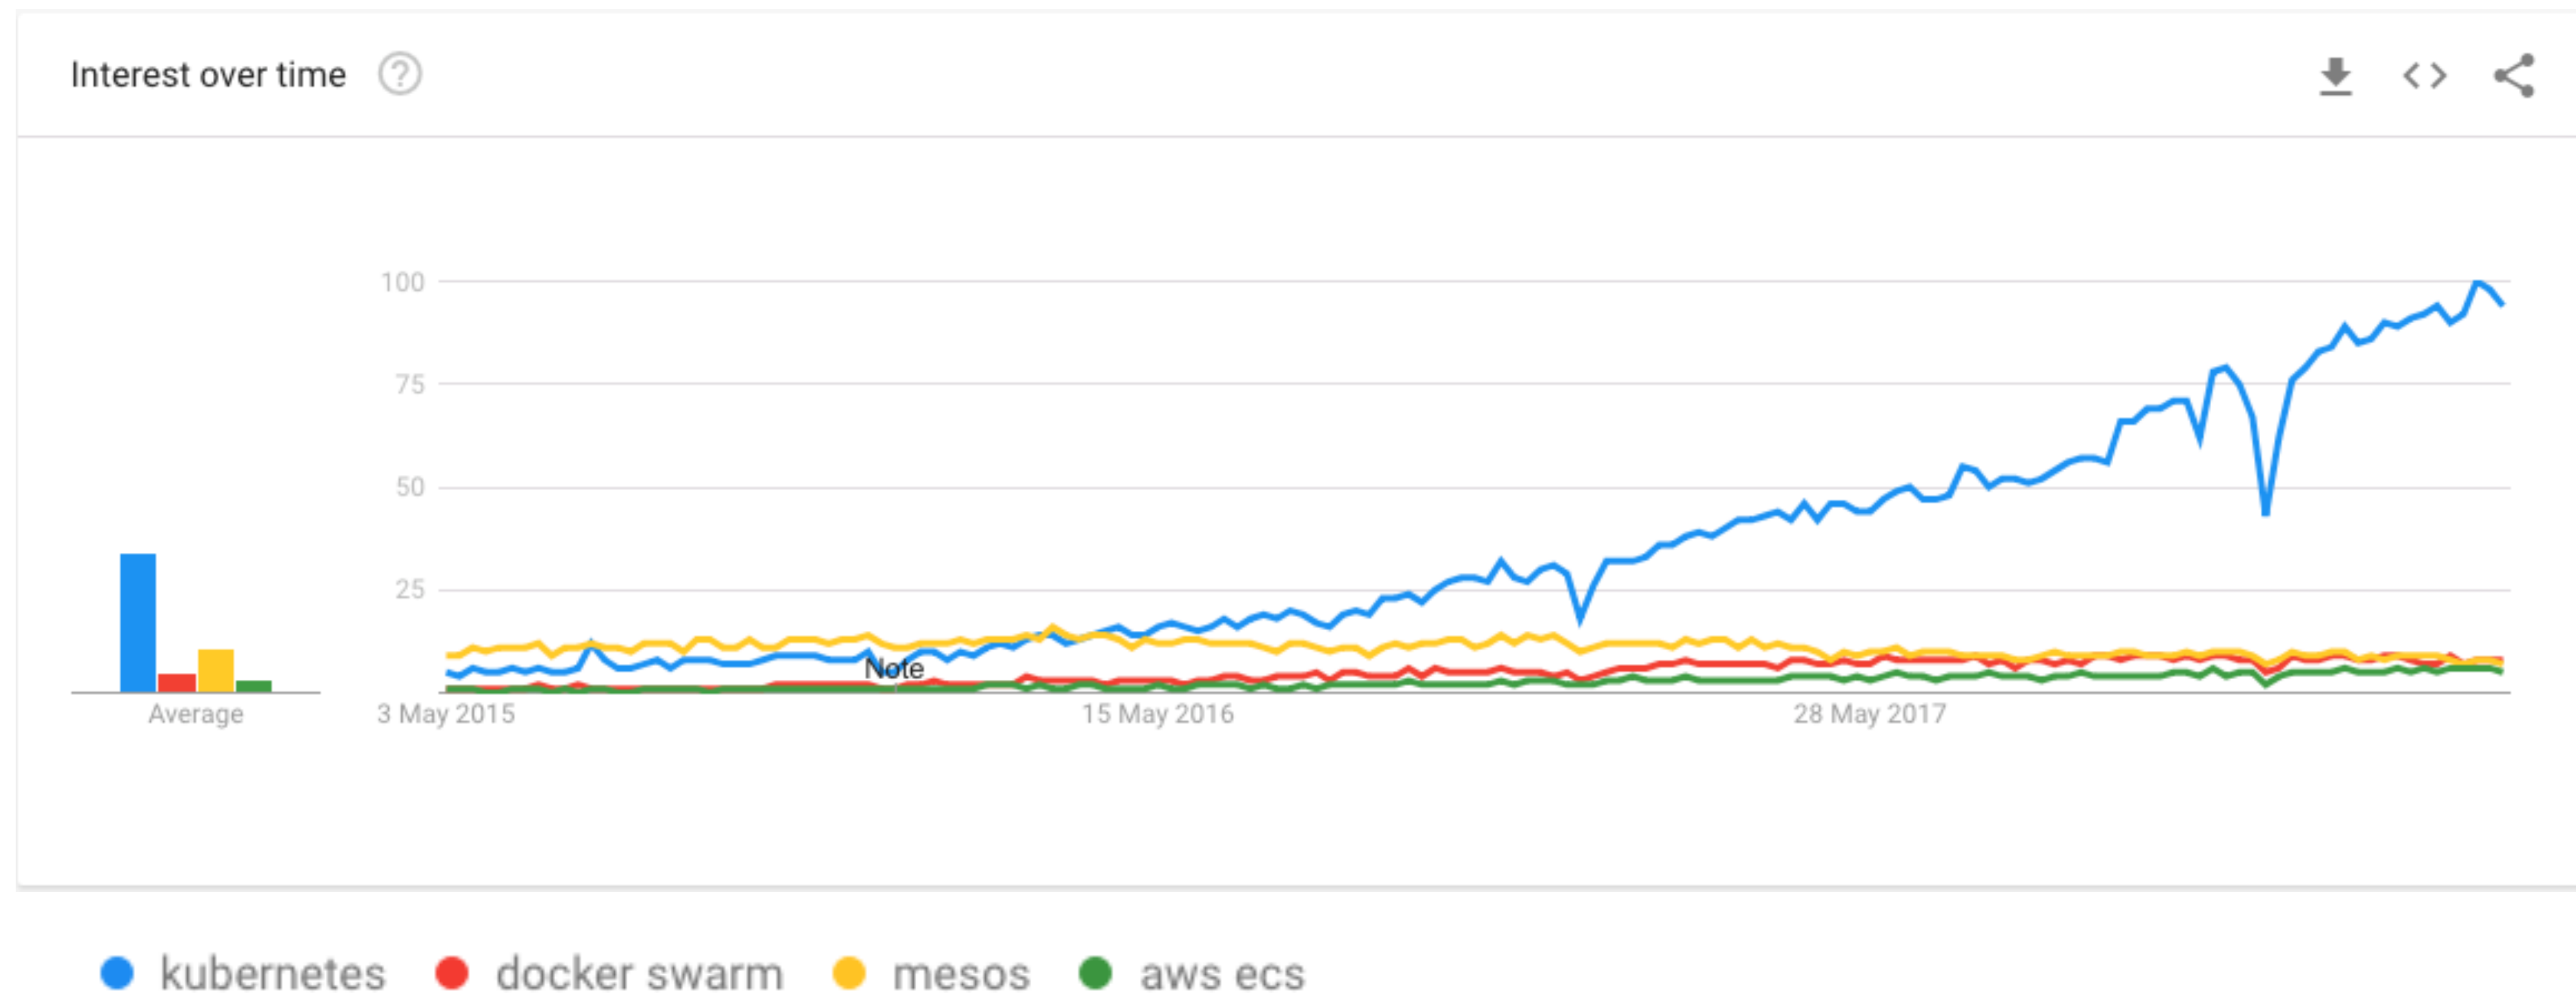
# Kubernetes

On-Prem or Cloud Agnostic Kubernetes

# Introvideo

# On-Prem / Cloud Agnostic Kubernetes

Kubernetes, the **most popular** orchestration tool!



# On-Prem / Cloud Agnostic Kubernetes

---

- Only start this course:
  - If you finished the Course “**Learn DevOps: The Complete Kubernetes Course**”
    - Which you can find on the Udemy marketplace
  - Or, if you are an experienced Kubernetes user
    - I will assume in this course that you know how to run apps on Kubernetes, and know the tooling around this

# Course Overview

Kubernetes topic	Technology
Installing kubernetes on-prem	kubeadm
File, Block, and Object storage	Kubernetes Operators, Rook with ceph
Managing SSL (HTTPS apps & endpoints)	cert-manager
LDAP authentication	Dex with LDAP
Service Mesh, Load Balancing and proxying	Envoy, Istio
Networking	Calico
Secret store	Vault
PaaS	OpenShift Origin

# Who am I

---

- My name is Edward Viaene
- I am a **consultant** and **trainer** in Cloud and Big Data technologies
- I'm a big advocate of **Agile** and **DevOps techniques** in all the projects I work on
- I held various roles from **banking** to **startups**
- I have a **background** in Unix/Linux, Networks, Security, Risk, and distributed computing
- Nowadays I specialize in everything that has to do with **Cloud** and **DevOps**

# On-Prem or Cloud Agnostic Kubernetes

Introduction

# On-Prem or Cloud Agnostic Kubernetes

---

- Only start this course:
  - If you finished the Course “**Learn DevOps: The Complete Kubernetes Course**”
    - Which you can find on the Udemy marketplace
  - Or, if you are an experienced Kubernetes user
    - I will assume in this course that you know how to run apps on Kubernetes, and know the tooling around this



# Course Overview

Kubernetes topic	Technology
Installing kubernetes on-prem	kubeadm
File, Block, and Object storage	Kubernetes Operators, Rook with ceph
Managing SSL (HTTPS apps & endpoints)	cert-manager
LDAP authentication	Dex with LDAP
Service Mesh, Load Balancing and proxying	Envoy, Istio
Networking	Calico
Secret store	Vault
PaaS	OpenShift Origin

# Course objectives

---

- To be able to **use** Kubernetes on-prem or in a Cloud Agnostic way
  - This allows you to **use** Kubernetes in an enterprise environment
- After this course you should be able to deploy Kubernetes anywhere
  - using your own integrations
  - like storage, certificates, authentication, and so on

# Who am I

---

- My name is Edward Viaene
- I am a **consultant** and **trainer** in Cloud and Big Data technologies
- I'm a big advocate of **Agile** and **DevOps techniques** in all the projects I work on
- I held various roles from **banking** to **startups**
- I have a **background** in Unix/Linux, Networks, Security, Risk, and distributed computing
- Nowadays I specialize in everything that has to do with **Cloud** and **DevOps**

# Online Training

---

- **Online training** on Udemy
  - **DevOps, Distributed Computing, Cloud, Terraform, Big Data (Hadoop)**
  - Using online video lectures
  - 45,000+ enrolled students in 100+ countries

# Kubernetes Introduction

Trail map

# Kubernetes installation

using kubeadm

# Kubeadm

---

- **Kubeadm** is a **toolkit** by Kubernetes to create a cluster
- It works on **any deb / rpm compatible Linux OS**, for example Ubuntu, Debian, Redhat, or CentOS
  - This is the main **advantage** of kubeadm, because **a lot of tools are OS / Cloud specific**
- The tool itself is still in beta (Q1 2018), but is expected to become **stable** somewhere **this year**
- It's very **easy to use** and lets you spin up your Kubernetes cluster in just a **couple of minutes**

# Kubeadm

---

- Kubeadm supports **bootstrap tokens**
  - Those are **simple tokens** that can be used to **create a cluster**, or to **join nodes** later on
  - The **tokens** are of the **format** *abcdef.0123456789abcdef*
- Kubeadm supports upgrading/downgrading clusters
- It **does not install a networking solution**
  - You have to install a **C**ontainer **N**etwork **I**nterface - compliant network solution yourself using `kubectrl apply` (as I will show in the demo)



# Kubeadm prerequisites

---

- deb / rpm compatible system (or CoreOS' Container Linux)
- **2 GB** of memory
- **2 CPUs** for the master node
- Network connectivity between the nodes
  - Can be a **private network** (internal IP addresses)
  - Or **public routable internet IP addresses** (in this case its best to use a firewall to only allow access within the cluster and to the users)
- Typically, you need minimal 2 nodes (one master node and one to schedule pods on)

# Demo

---

- In the demo I will use DigitalOcean to spin up droplets (VMs)
- You can get \$10 worth of credits if you use the following link to sign up:
  - <https://m.do.co/c/007f99ffb902>
- A **2 GB memory droplet currently costs \$10 / month**, but it is **billed per hour**, so you can run 2x 2GB RAM droplets for half a month or 4x 2GB RAM droplets for bit more than a week

# Demo

installing Kubernetes using kubeadm

# Operators

# Operators

---

- An Operator is a method of **packaging, deploying, and managing** a Kubernetes Application (definition: <https://coreos.com/operators/>)
- It puts **operational knowledge** in an application
  - It brings the user **closer to the experience of managed cloud services**, rather than having to know all the specifics of an application deployed to Kubernetes
  - Once an Operator is deployed, it can be **managed using Custom Resource Definitions** (arbitrary types that extend the Kubernetes API)
- It also provides a great way to deploy Stateful services on Kubernetes (because a lot of complexities can be hidden from the end-user)

# Custom Resource Definitions

---

- Custom Resource Definitions (CRDs) are **extensions of the Kubernetes API**
- It allows the Kubernetes user to use **custom objects** (the objects you use in yaml files), and create / modify / delete those objects on the cluster
  - For example: you could run a *kubectl create* on a yaml file containing a custom database object, to spin up a database on your cluster
- The custom objects are **not necessarily available on all clusters**
  - They can be dynamically registered / deregistered
- **Operators include CRDs**
  - By adding an Operator, you'll register these custom resource definitions

# Operators example

---

- **etcd, Rook, Prometheus, and Vault** are examples of technologies can be deployed as an Operator
- Let's use **etcd** as an example (etcd is a distributed key value store)
  - Once the etcd operator is deployed, a new etcd cluster can be created by using the following yaml file:

```
apiVersion: "etcd.database.coreos.com/v1beta2"
kind: "EtcdCluster"
metadata:
  name: "example-etcd-cluster"
spec:
  size: 3
  version: "3.2.13"
```

# Operators example

---

- Resizing the cluster is now just a matter of changing the yaml file:

```
apiVersion: "etcd.database.coreos.com/v1beta2"
kind: "EtcdCluster"
metadata:
  name: "example-etcd-cluster"
spec:
  size: 5 # was 3 previously
  version: "3.2.13"
```

- After making the edit, the changes can be applied using *kubectl apply*
- The same is true for **the version** number: to upgrade the etcd cluster you just change the version, enter *kubectl apply*, and the etcd cluster will be upgraded



# Operators

---

- Using operators **simplifies deployment and management** a lot
- This example was using an etcd cluster, but more software is being released using operators for Kubernetes
  - For example the PostgreSQL operator by Zalando (<https://github.com/zalando-incubator/postgres-operator>)
  - Or the MySQL Operator, providing you with a simple API to create a MySQL Database (<https://github.com/oracle/mysql-operator>):

```
apiVersion: mysql.oracle.com/v1
kind: MySQLCluster
metadata:
  name: myappdb
```

# Operators

---

- You can also build your own operators, using the following tools:  
(Source: <https://coreos.com/operators/>)
  - **The Operator SDK:** makes it easy to build an operator, rather than having to learn the Kubernetes API specifics
  - **Operator Lifecycle Manager:** oversees installation, updates, and management of the lifecycle of all the operators
  - **Operator Metering:** Usage reporting
- I'll be using **Operators** in this course
  - The next lectures will be about **Rook**, which will be deployed using Operators

# Rook

Introduction to Rook

# Rook

---

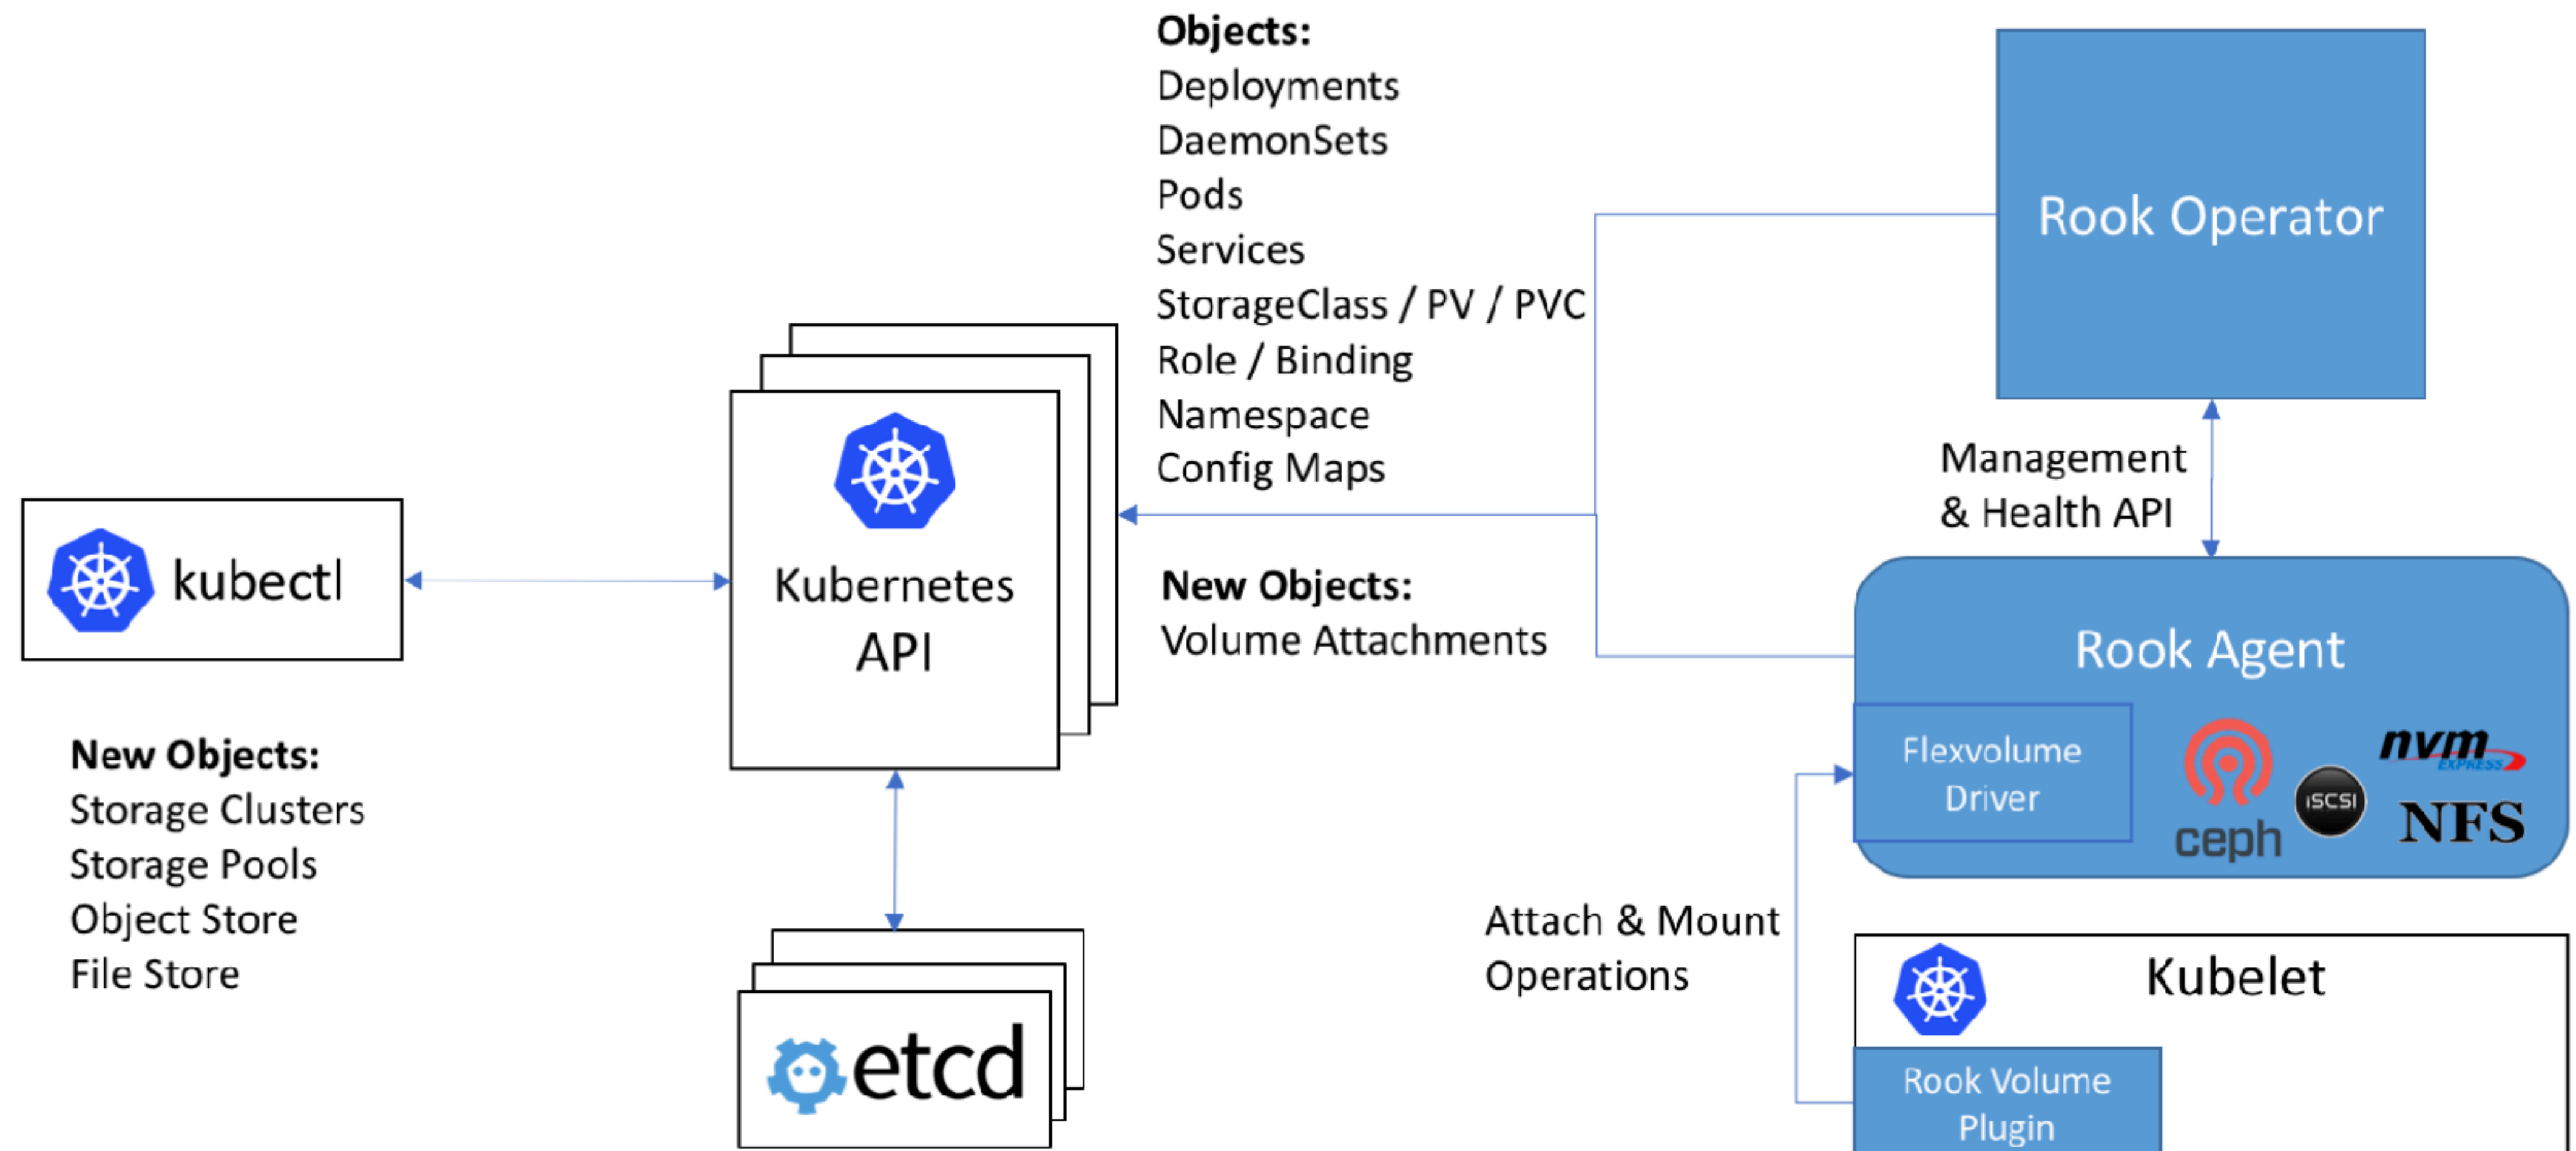
- Rook is an **open source orchestrator** for **distributed storage systems** running in Kubernetes. (definition: <https://rook.io/docs/rook/master/>)
  - Rook allows you to **use storage systems on Kubernetes clusters** (that cannot use public cloud storage, or want to be cloud agnostic)
  - If you're on the public cloud, it's very easy to **attach a storage volume to a pod**, to allow your app to persist its data, even when the pod or node shuts down
  - It's not that easy when you're **not on the major cloud providers** like AWS / Azure / Google Cloud
- **Rook** wants to make it easy for you to use a **storage system**, even when you're not on one of those major cloud providers, or using an **on-prem cluster**

# Rook

---

- Rook **automates** the **configuration, deployment, maintenance** of distributed storage software
- This way, you **don't need to worry** about the **difficulties of setting up storage systems**
  - Rook will **orchestrate** all this management for you
- Rook is currently (early 2018) in alpha, but Rook already looks very promising and will sure be stable at some point soon
- Currently Rook uses **Ceph** as underlying storage, but **Minio** and **CockroachDB** are also available
  - **More storage engines** will be added in future releases

# Rook



Source: <https://rook.io/docs/rook/master/>

Ceph



# Ceph

---

- Ceph is software that provides **object, file, and block storage**
- It's **open source**
- It's **distributed without a single point of failure**
- Ceph **replicates** its data to make it **fault tolerant** (a node can fail, and you still have your data available)
- It's **self-healing** and **self-managing**
- Scalable to exabyte level



# Ceph

---

- Ceph provides **3 different types of storage**:
  - **File Storage**: to **store files and directories**, similar to accessing files over Networking File System (NFS), or using a Network Attached Storage (NAS), or EFS (Elastic File System) on AWS
  - **Block Storage**: like a hard drive, to store data using a filesystem. A database needs block storage. Examples are a SAN (Storage Area Network, which can provide block storage to servers), or EBS (Elastic Block Storage) on AWS
    - Typical use case is to store files for your OS, storage for databases, etc
  - **Object Storage**: To store any type of data as an **object**, identified by a **key**, with the possibility to add **metadata**. This type of storage lends itself to be **distributed** and **scalable**. For example, AWS S3 provides Object Storage
    - Can be used to store **unstructured data** like pictures, website assets, videos, log files, etc

# Ceph Components

---

- Ceph has multiple components:
  - **Ceph Monitor** (min 3): **maintains a map of the cluster state** for the other ceph components (the daemons) to communicate. Also responsible for **authentication** between daemons and clients
  - **Ceph Manager daemon**: responsible to **keep track of runtime metrics** and the cluster state
  - **Ceph OSDs** (Object Storage Daemon, min 3): **stores the data**, is responsible for **replication, recovery, rebalancing** and provides information for the monitoring daemons
  - MDSs (**Ceph Metadata Server**): stores metadata for the **Ceph FileSystem** storage type (not for block/object storage types)
- Ceph stores data as **objects** within **logical storage pools**
- It uses the **CRUSH algorithm** (Controlled Replication Under Scalable Hashing), which allows Ceph to be scalable

# Ceph Components

---

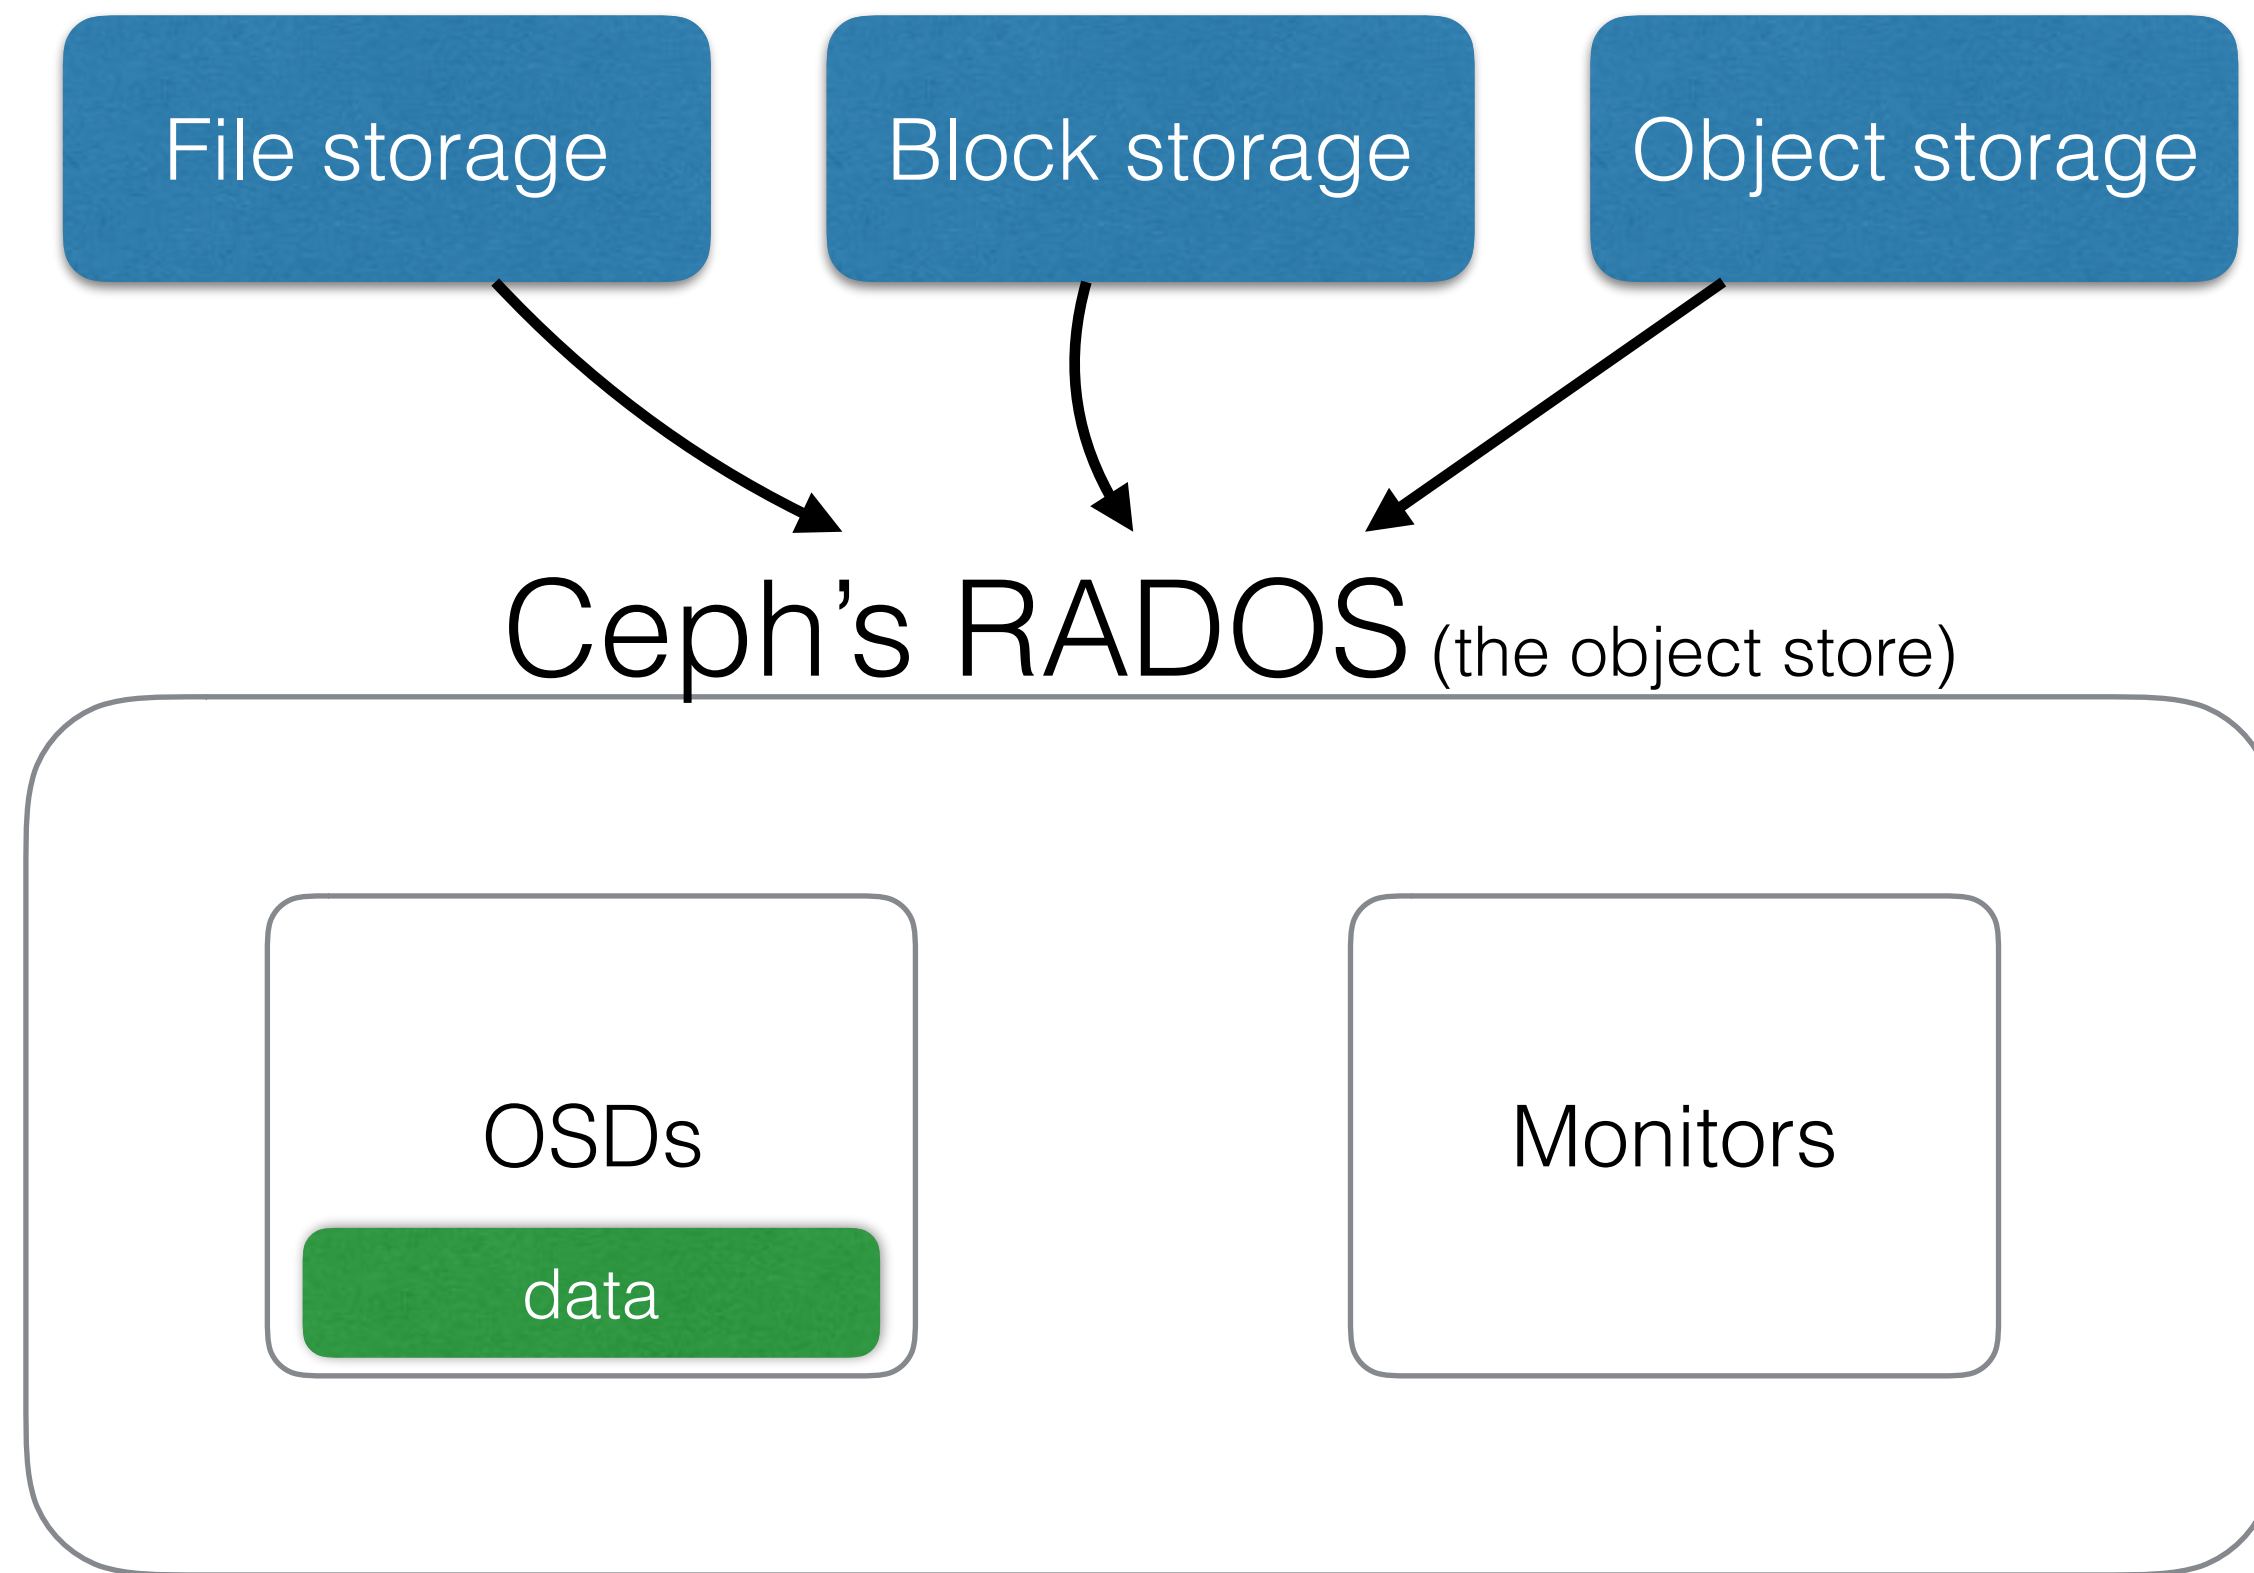
- Object storage in Ceph is provided by the **distributed object storage mechanism within Ceph**
- Ceph software libraries (librados) provides clients access to the “Reliable Autonomic Distributed Object Store” (RADOS)
  - RADOS provides a **reliable, autonomous, distributed object store** comprised of **self-healing, self-managing, intelligent storage** nodes  
(definition source: <http://docs.ceph.com/docs/master/architecture/>)
  - There is also a RESTful interface that can provide an AWS S3 compatible interface to this object store

# Ceph Components

---

- Ceph **block storage** is provided by Ceph's RADOS Block Device (RBD)
- RBD is built on top of the same **Ceph Object Storage**
  - Ceph stores **the block images** as **objects** in the object store
  - It's also built on librados, the software library

# Ceph Components



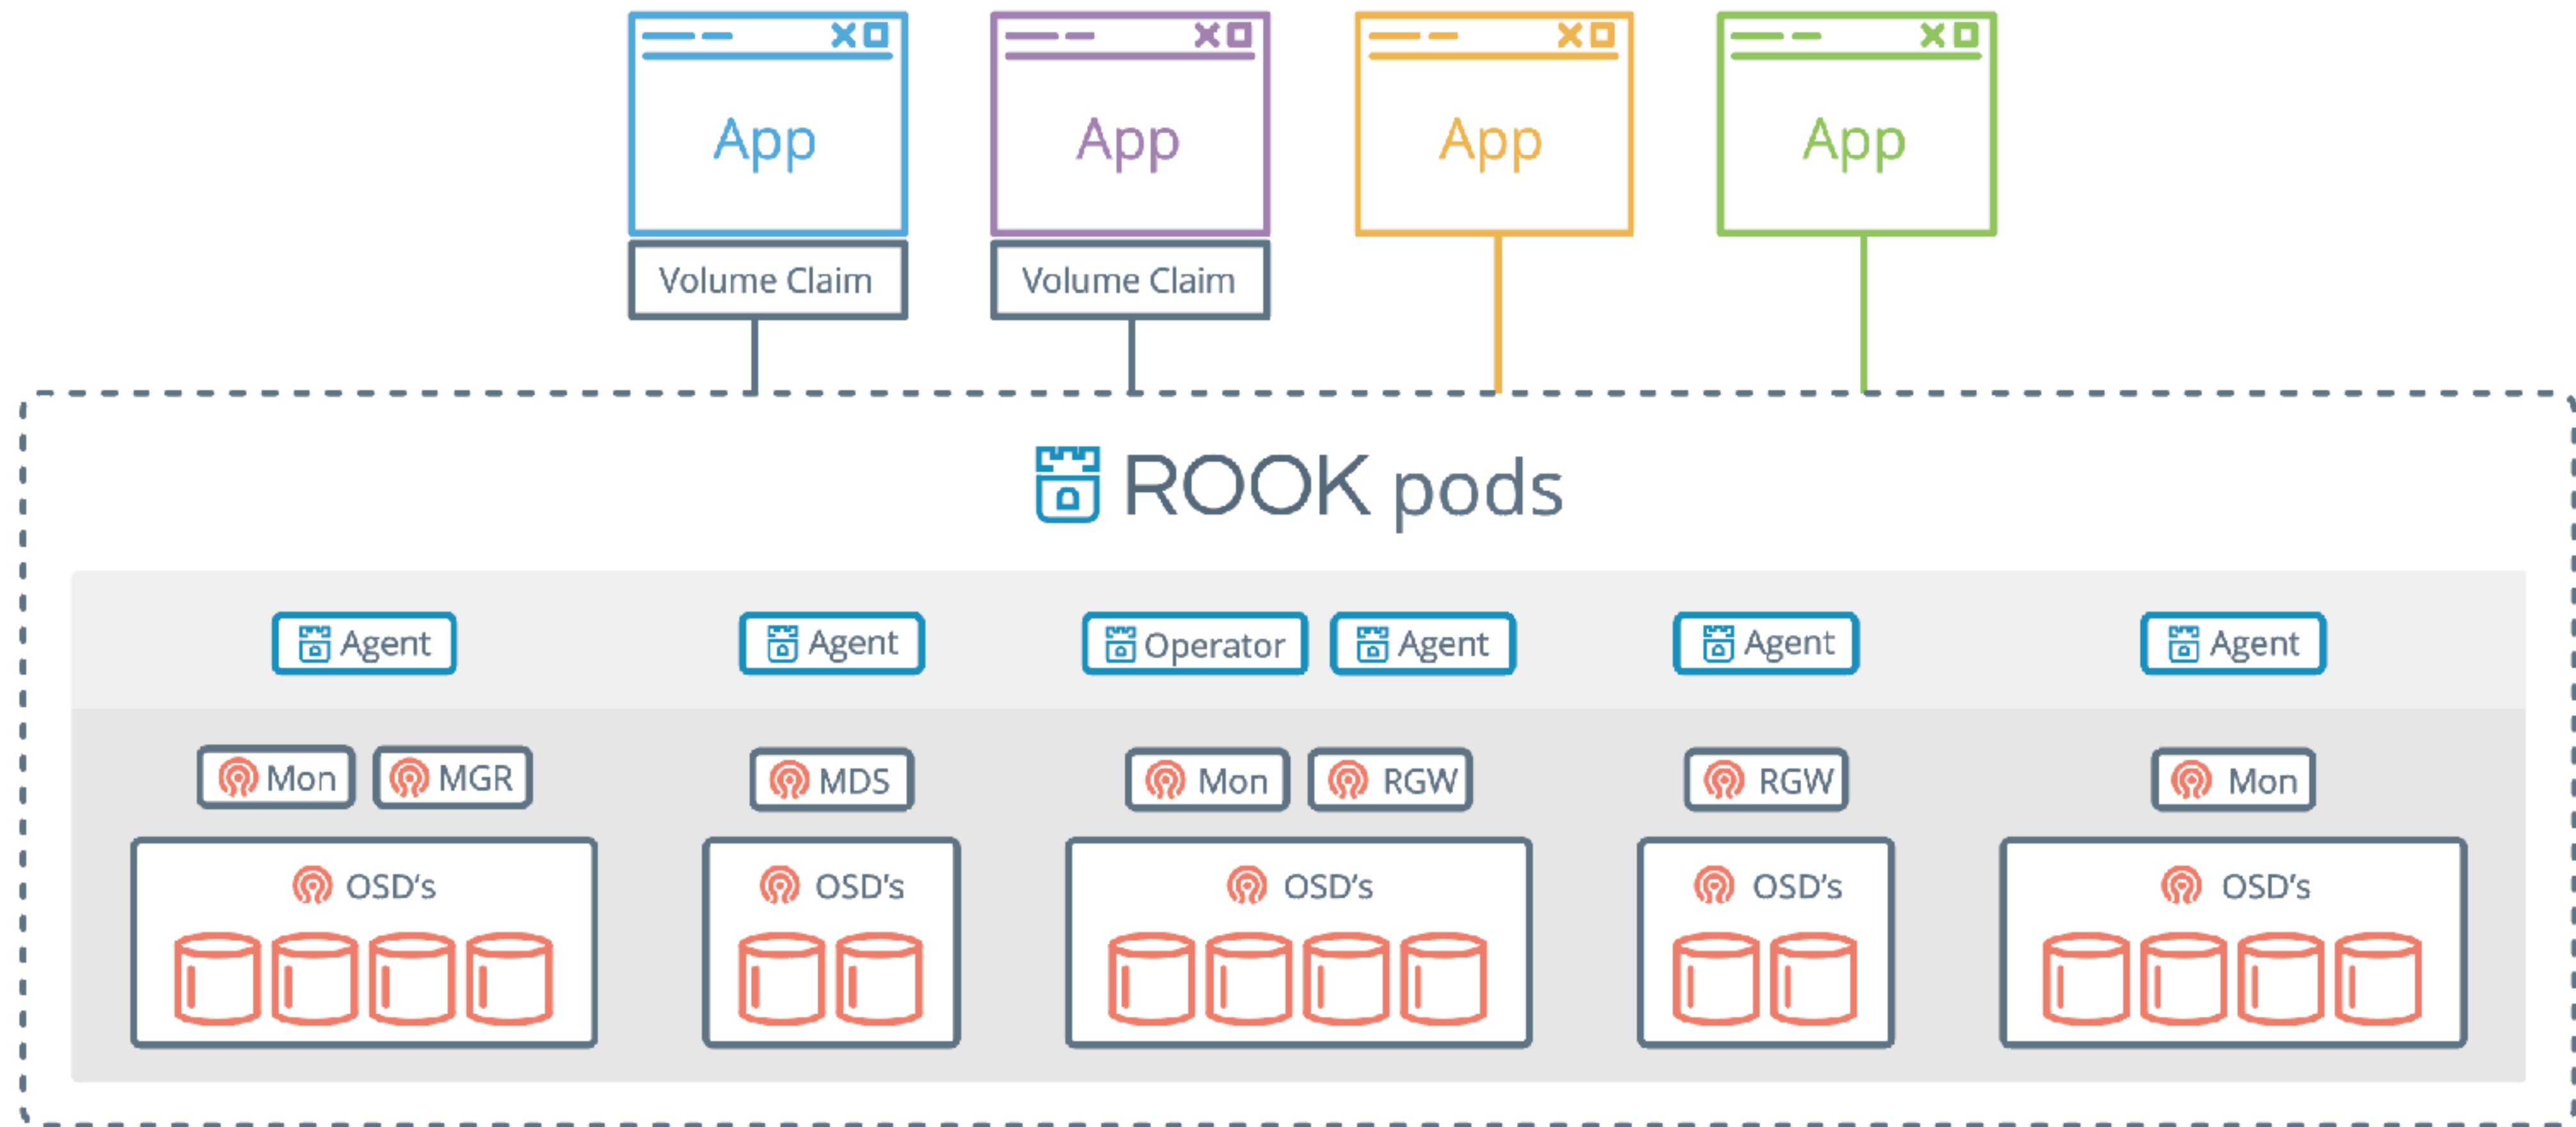
- Data can come in from Ceph's **file storage**, **block storage**, or **object storage**, and Ceph will **store this data as an object in Ceph**
- Each object is stored within the Object Storage Device (OSD)
- The OSDs will run on multiple nodes
- They will handle the read/write operations to their underlying storage

Ceph with Rook



# Ceph on Rook

## Rook Architecture



Source: <https://rook.io/docs/rook/master/>

# Ceph on Rook

---

- Rook supports all 3 types of storage: **block, file and object storage**
- I will use in the demo **Ceph** for all these types, but **other backends** are also a possibility
  - Rook will do a good job to **abstract this away from you**, so most of the configuration is nicely hidden from you
  - You 'll be able to use the **kubernetes yaml files** to set configuration options



# Deployment steps

---

- First, you'll need to deploy the **rook operator**
  - Using the provided yaml files
  - Using the helm chart
- Then, you can **create the rook cluster**
  - Also using yaml definitions (this time using: apiVersion: rook.io/v1alpha1)
  - This will use the rook operator rather than the Kubernetes API
- After that, **block / file / object storage can be configured**
  - Using the rook API and Kubernetes storage API - using this storage API means using rook storage will become as easy as using for example AWS EBS or NFS

# Demo

Ceph with Rook

# Demo

Rook Object Storage

# Demo

Rook Shared File System

cert-manager

# Cert-manager

---

- If you want to use a **secure http connection** (https), you need to have **certificates**
- Those certificates can be **bought**, or can be **issued** by some **public cloud providers**, like AWS's Certificate Manager
- Managing SSL / TLS certificates yourself often **takes a lot of time** and are **time consuming to install and extend**
  - You also **cannot issue your own certificates** for production websites, as they are not trusted by the common internet browsers (Chrome, IE, ...)
- **Cert-manager** can **ease** the **issuing** of certificates and the **management** of it

# Cert-manager

---

- Cert-manager can use letsencrypt
- Let's encrypt is a **free**, **automated** and **open** Certificate Authority  
(definition: <https://letsencrypt.org/>)
  - Let's encrypt can issue certificates for free for your app or website
  - You'll need to prove to *let's encrypt* that you are the owner of a domain
  - After that, they'll issue a certificate for you
  - The certificate is **recognized by major software vendors and browsers**

# Cert-manager

---

- Cert-manager can **automate the verification process** for let's encrypt
- With *Let's encrypt* you'll also have to **renew certificates** every **couple of months**
- Cert-Manager will **periodically check the validity** of the certificates and will **start the renewal process** if necessary
- *Let's encrypt* in combination with cert-manager **takes away a lot of hassle** to deal with certificates, allowing you to **secure your endpoints** in an easy, affordable way

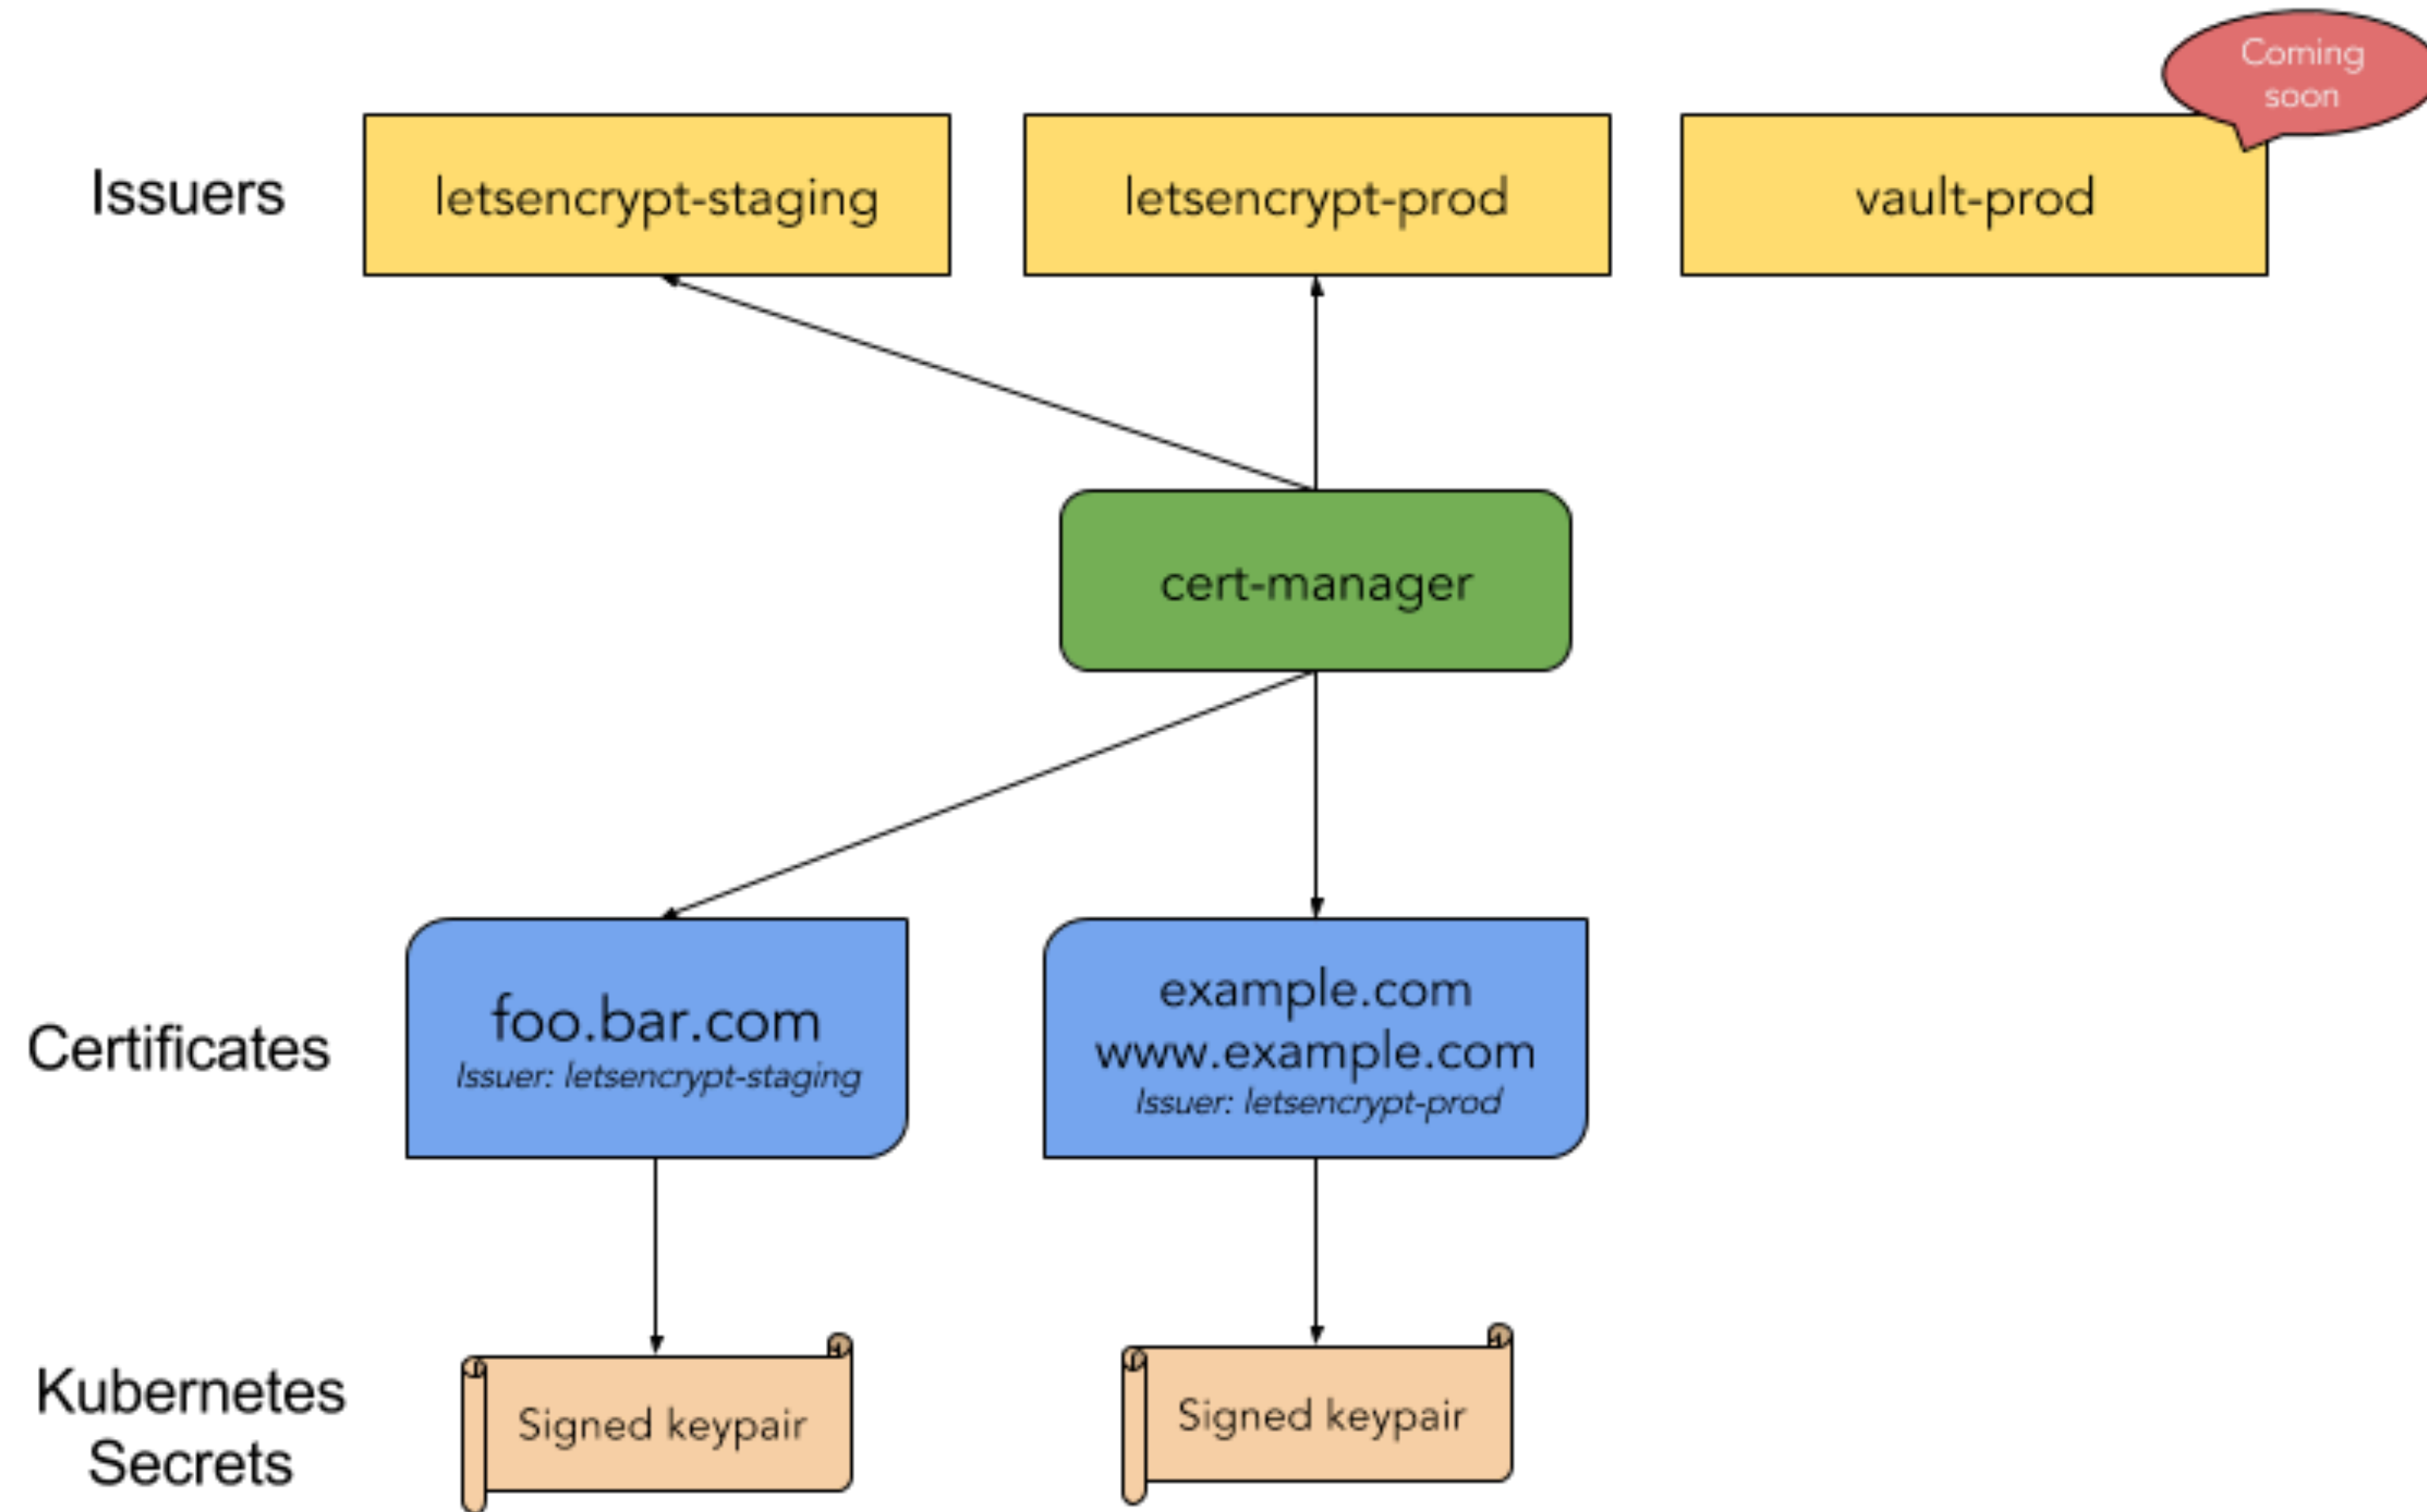


# Cert-manager

---

- You can only issue certificates for a **domain name you own**
- You'll need to have a domain name like xyz.com
  - If you were using a domain name to bring up your cluster for my first Kubernetes course, the “**Complete Kubernetes Course**”, you can **re-use this domain**
  - Otherwise, you can **get one for free** from www.dot.tk or other providers
  - Or, you can **buy one** through namecheap.com / AWS route53 / any other provider that sells domain names
  - Less popular extensions only cost **a few dollars**

# Cert-manager



From: <https://cert-manager.readthedocs.io/en/latest/index.html>

# Demo

cert-manager

Dex


# Dex

---

- Dex is **Identity Service**
  - It uses **OpenID Connect** (OIDC)
- **Kubernetes** can use **Dex** to **authenticate** its **users** (using OIDC)
- Dex uses **connectors to authenticate a user** using another Identity Provider
  - This allows you to use Dex to authenticate users in Kubernetes using LDAP, SAML, GitHub, Microsoft, and others

# Dex Architecture

---

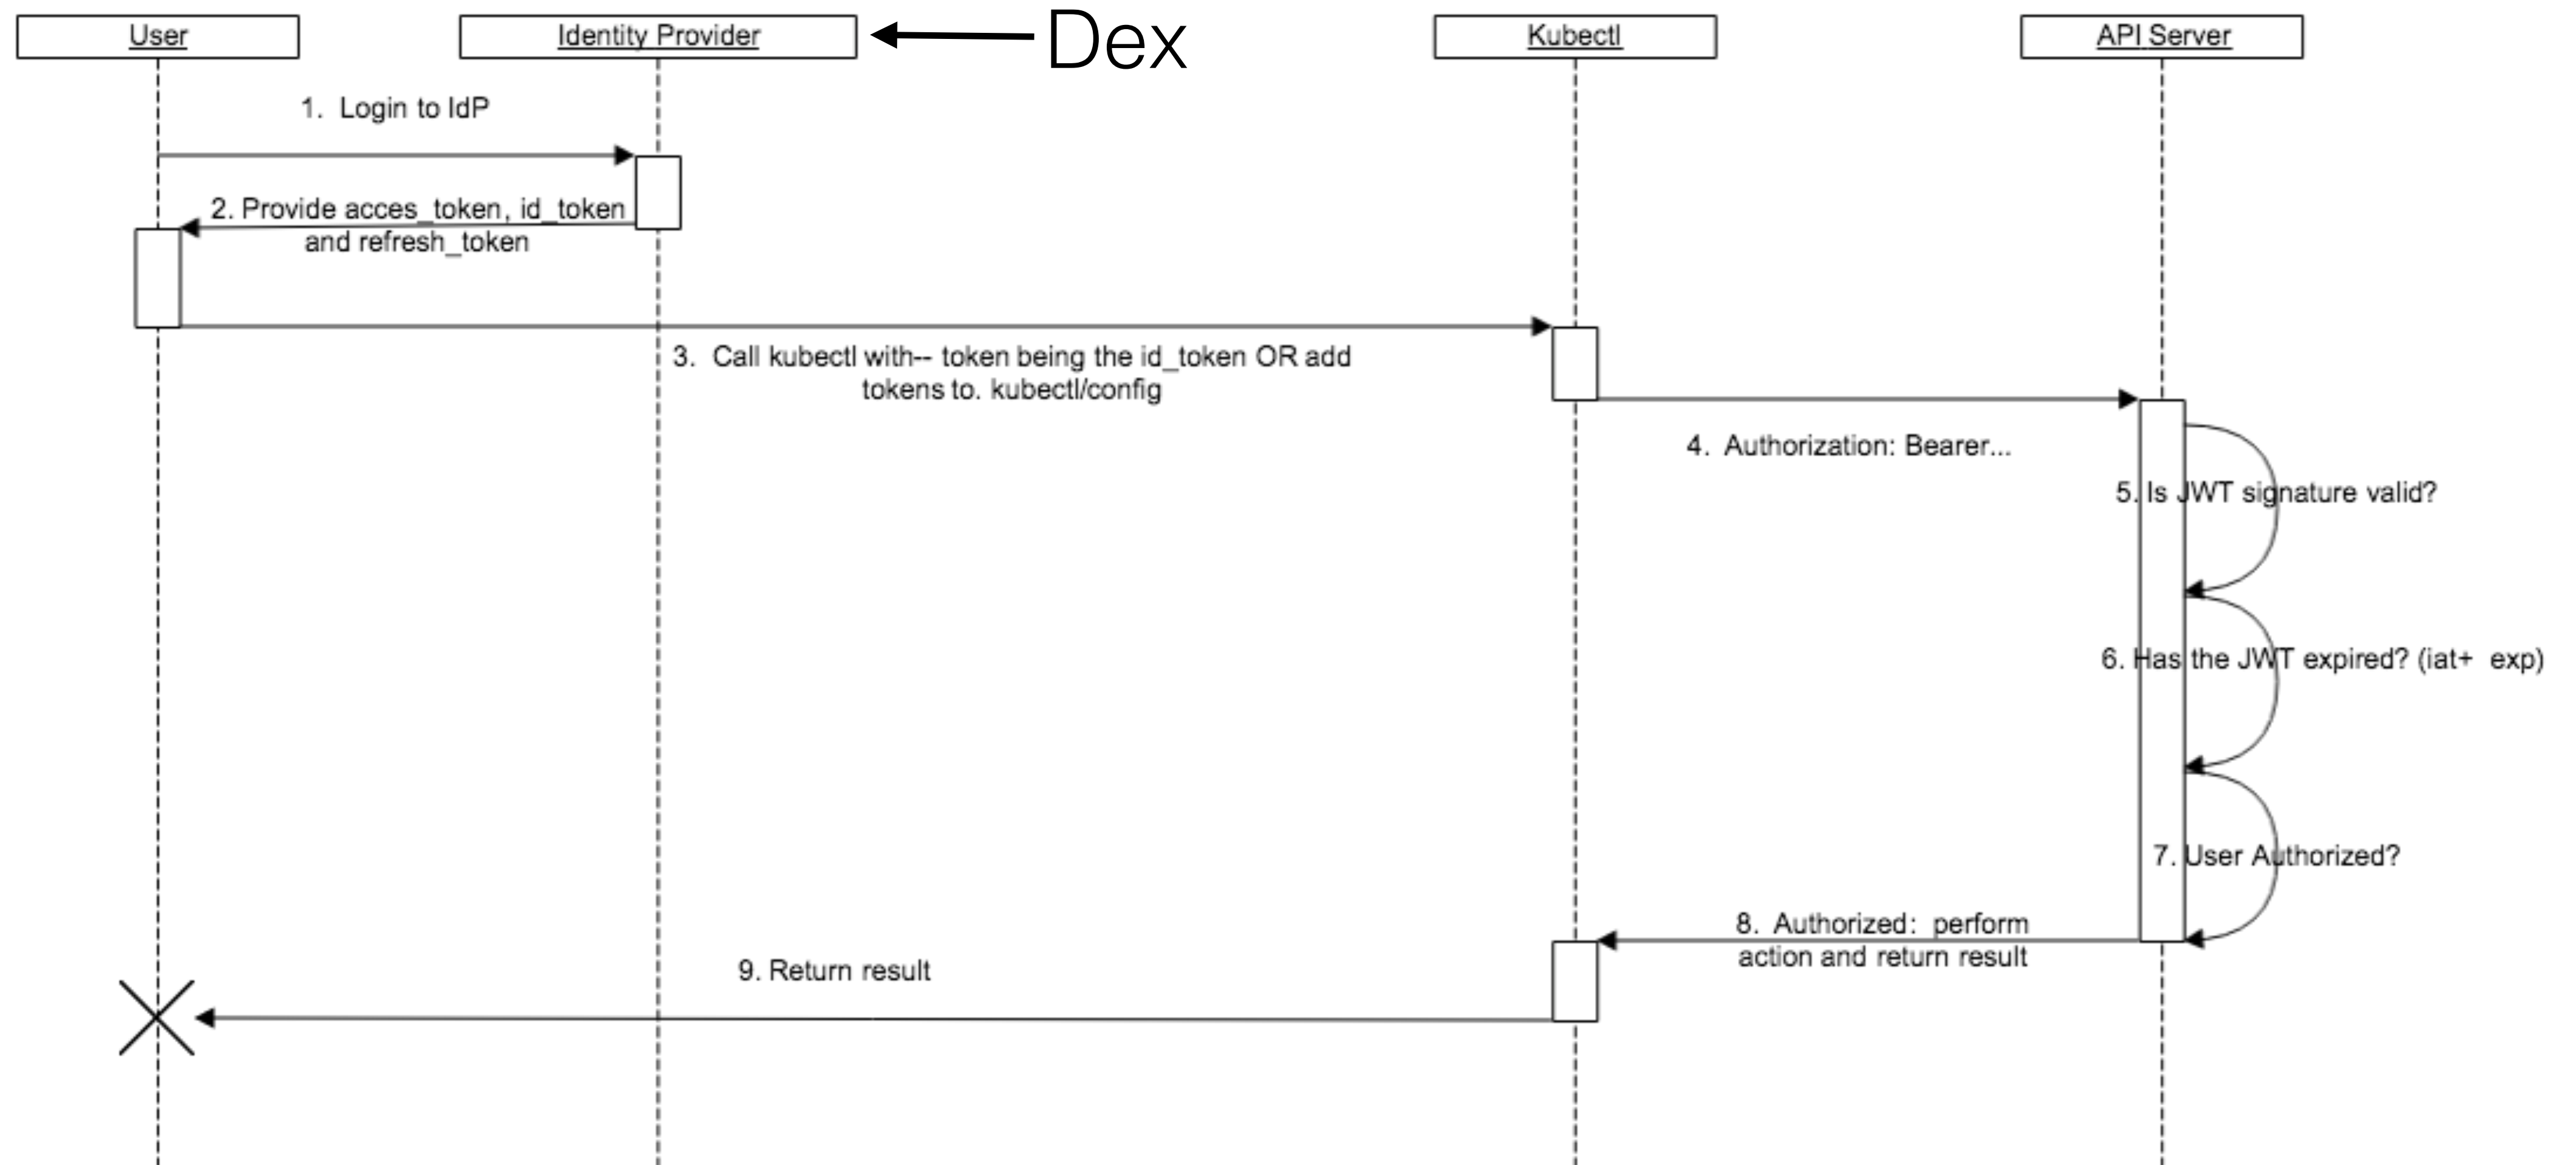


# Dex

---

- Most companies already have a **user directory**, using **OpenLDAP**, Microsoft **Active Directory** (which is LDAP compatible), or similar products
  - LDAP stands for **Lightweight Directory Access Protocol**
- It's less common for companies to already have an OpenID Connect (OIDC) implementation that you can use
- That's why you have to use software like Dex, that will **act as a bridge** between what **enterprises offer for authentication**, and what Kubernetes can use today
- Dex can use LDAP, but there are also **other connectors** you could use if your company doesn't use LDAP

# Kubernetes OIDC



Source: <https://kubernetes.io/docs/reference/access-authn-authz/authentication/>



# Demo

dex

# Demo

dex - LDAP

# Istio - Envoy

# Envoy

---

- When you break up a **monolith application** (1 codebase), into **micro-services** (multiple codebases), you end up with **lots of services** that need to be able to **communicate** with each other
- These communications between services need to be able to be **fast, reliable and flexible**
- To be able to implement this, you need a **service mesh**
  - A **service mesh** is an **infrastructure layer** for handling these service-to-service communications
  - This is usually implemented using proxies
  - Proxies manage these communications and ensure they're **fast, reliable and flexible**

# Envoy

---

- **Envoy** is a such a **proxy**
  - It is designed for **cloud native applications**
- Was originally built at **Lyft**
- Envoy is a **High Performance distributed proxy** written in C++
- You can see it as an **iteration** of the NGINX, HAProxy, hardware / cloud load balancers
- It's comparable with **Linkerd** (which is explained in my Advanced Kubernetes course)
  - While there's a lot of **overlap**, each solution has its **own distinct features**

# Envoy Features

---

- **Small memory** footprint
- HTTP/2 and gRPC support
  - It's a **transparent** HTTP/1.1 to HTTP/2 proxy
    - Not all browsers support HTTP/2 yet, so **incoming requests** can be HTTP/1.1, but **internally requests can be HTTP/2**
- **Advanced Loadbalancer Features** (automatic retries, circuit braking, rate limiting, request shadowing, zone load balancing, ...)
- **Configuration** can be **dynamically managed** using an **API**
- Native support for **distributed tracing**

# Comparison to linkerd

---

- **Linkerd** has **more features**, but that comes at a price of higher cpu and memory footprint
  - Linkerd is built on top of **Netty and Finagle** (JVM based), whereas **Envoy is written in C++**
  - If you're looking for **more features**, you might want to look at Linkerd, if you're looking for speed and low resource utilization, **Envoy wins**
    - **Istio**, discussed next, can give you the best of both worlds
- Linkerd **integrates** with **Consul and Zookeeper** for service discovery
- Envoy **supports hot reloading** using an API, Linkerd does not (by design)

Istio

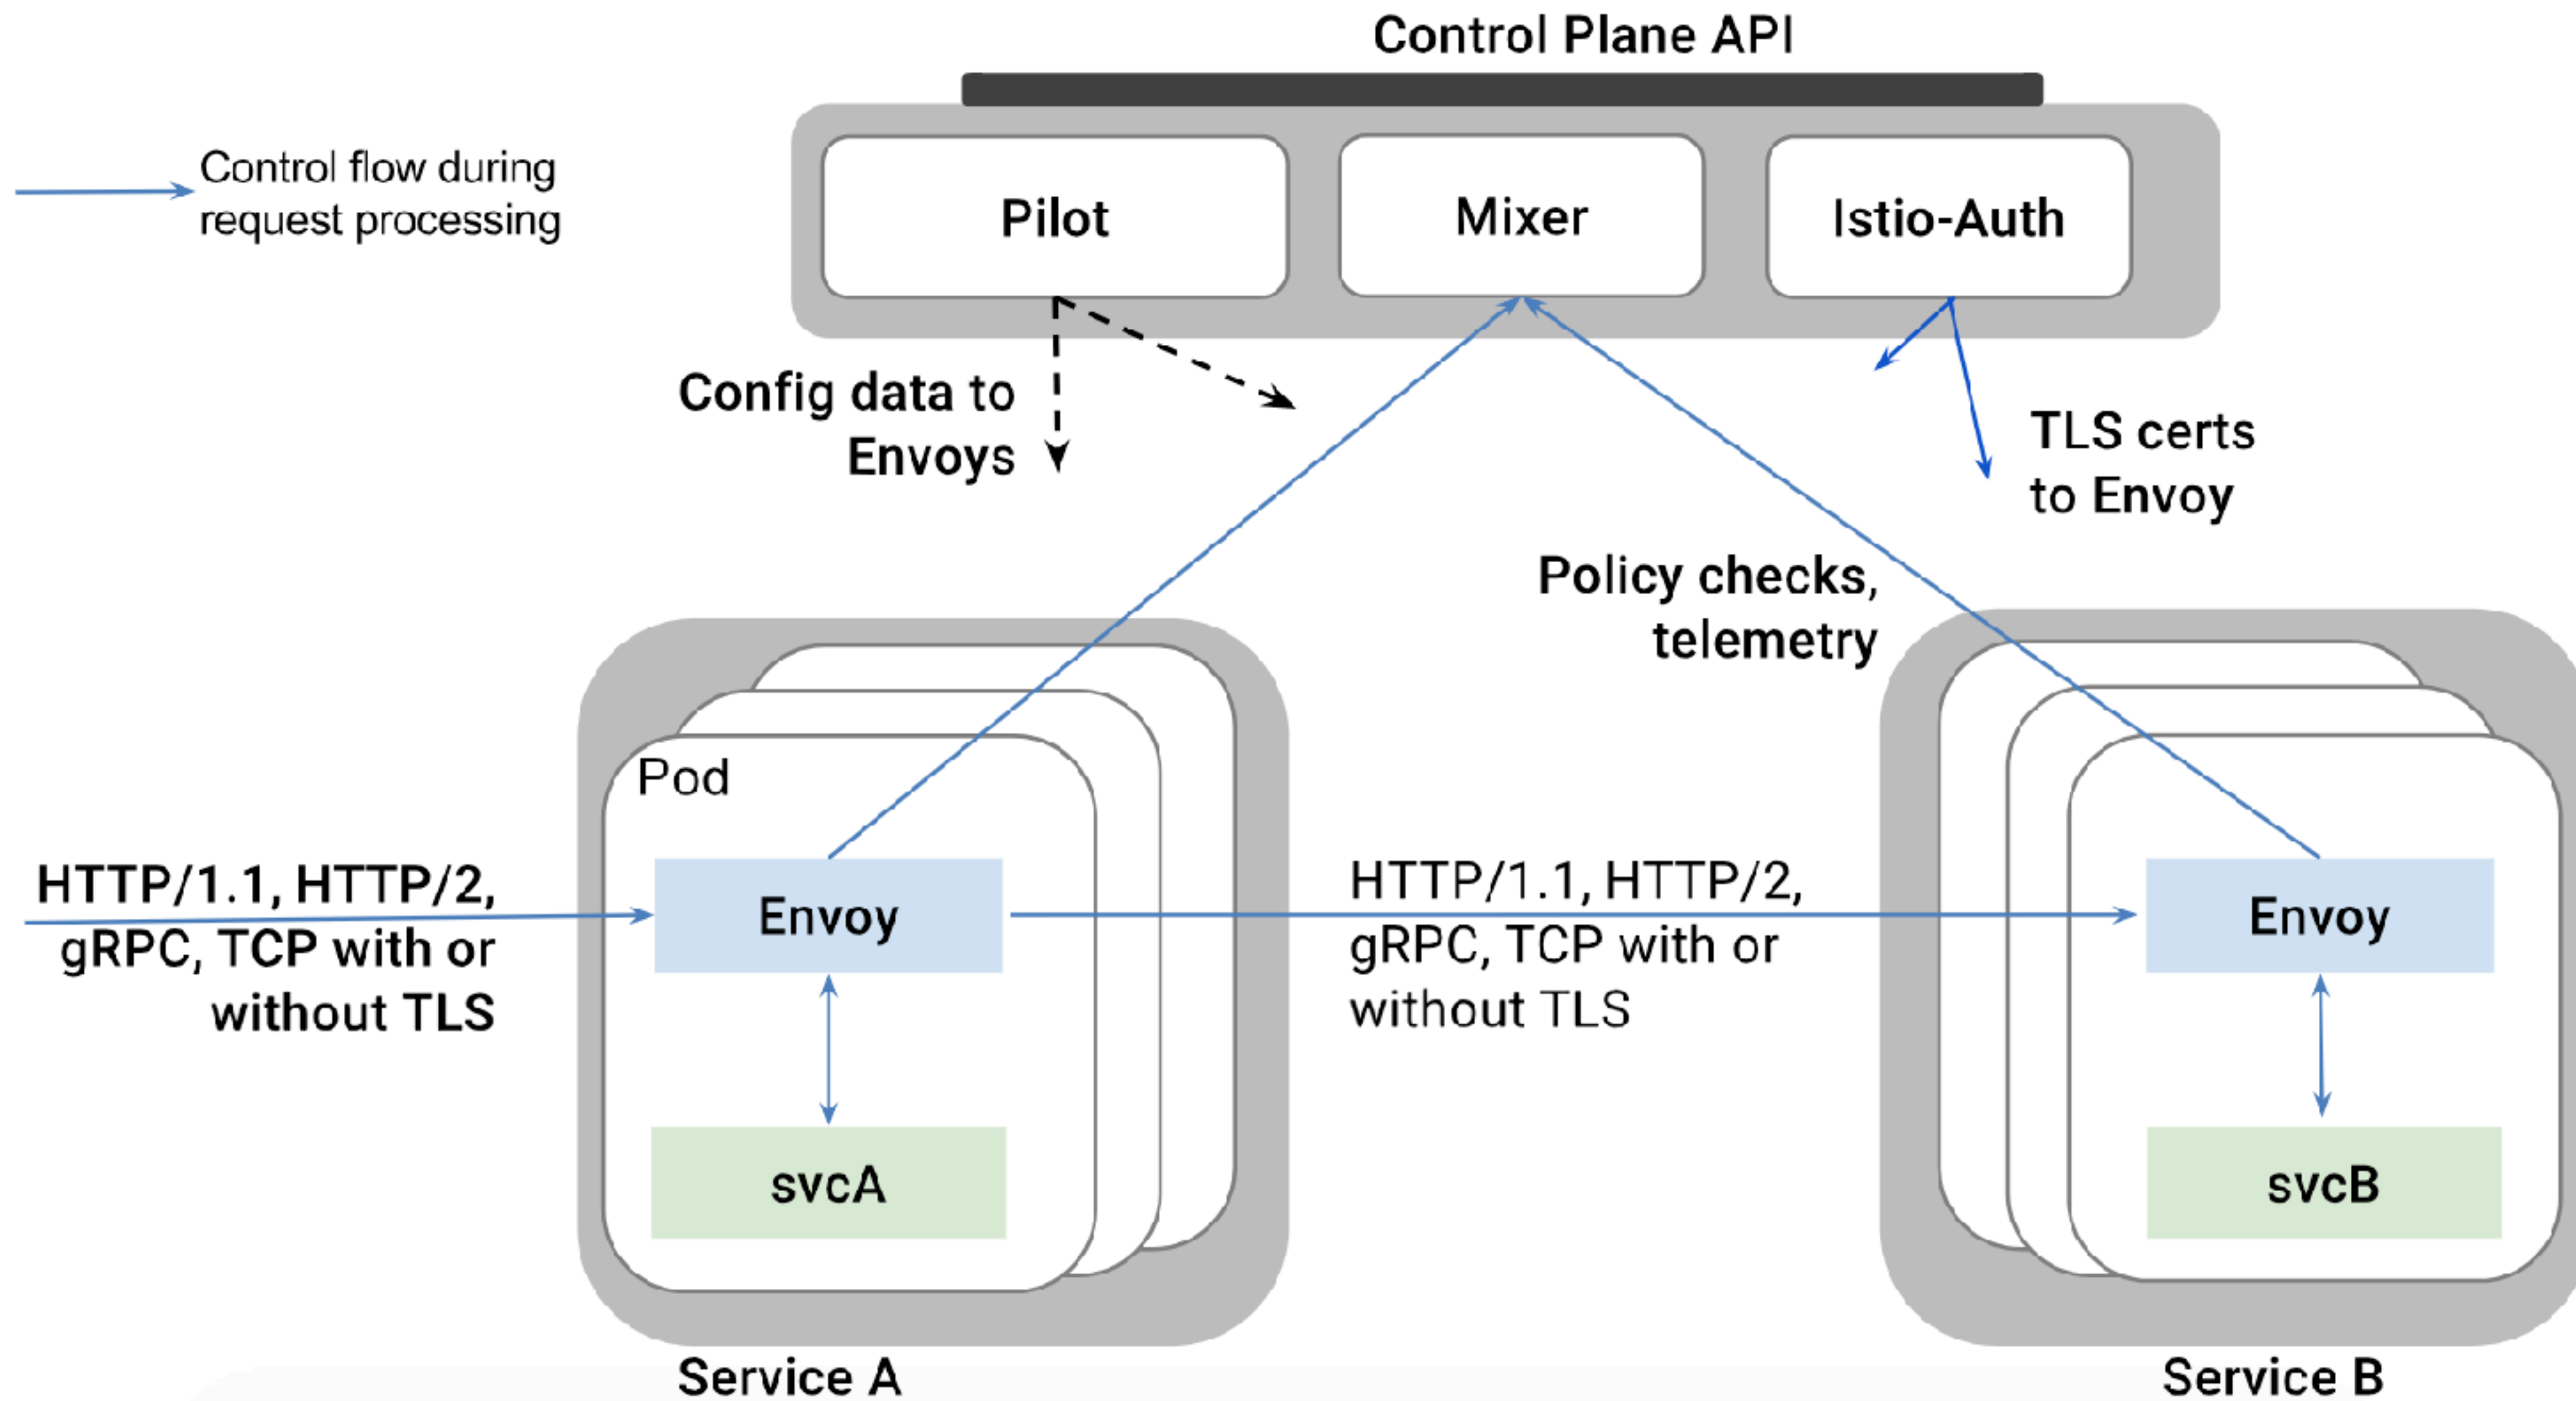


# Istio

---

- Istio is an **open platform** to **connect, manage, and secure microservices** (Definition: <https://istio.io/docs/concepts/what-is-istio/overview.html>)
- Key capabilities include:
  - It supports Kubernetes
  - Can **control traffic between** services, can make it more **robust** and **reliable**
  - Can show you **dependencies** and the **flow** between services
  - Provides **access policies** and **authentication** within your service mesh

# Istio Architecture



from: <https://istio.io/docs/concepts/what-is-istio/overview.html>)

# Istio Components

---

- Envoy (data plane)
  - Istio uses the **Envoy proxy** in its **data plane**
  - It uses a **sidecar deployment**, which means a deployment along the application (a one to one relation between app/pod and proxy)
- Mixer (control plane)
  - Responsible for enforcing **access control** and **usage policies**
  - Collects **telemetry** data from Envoy

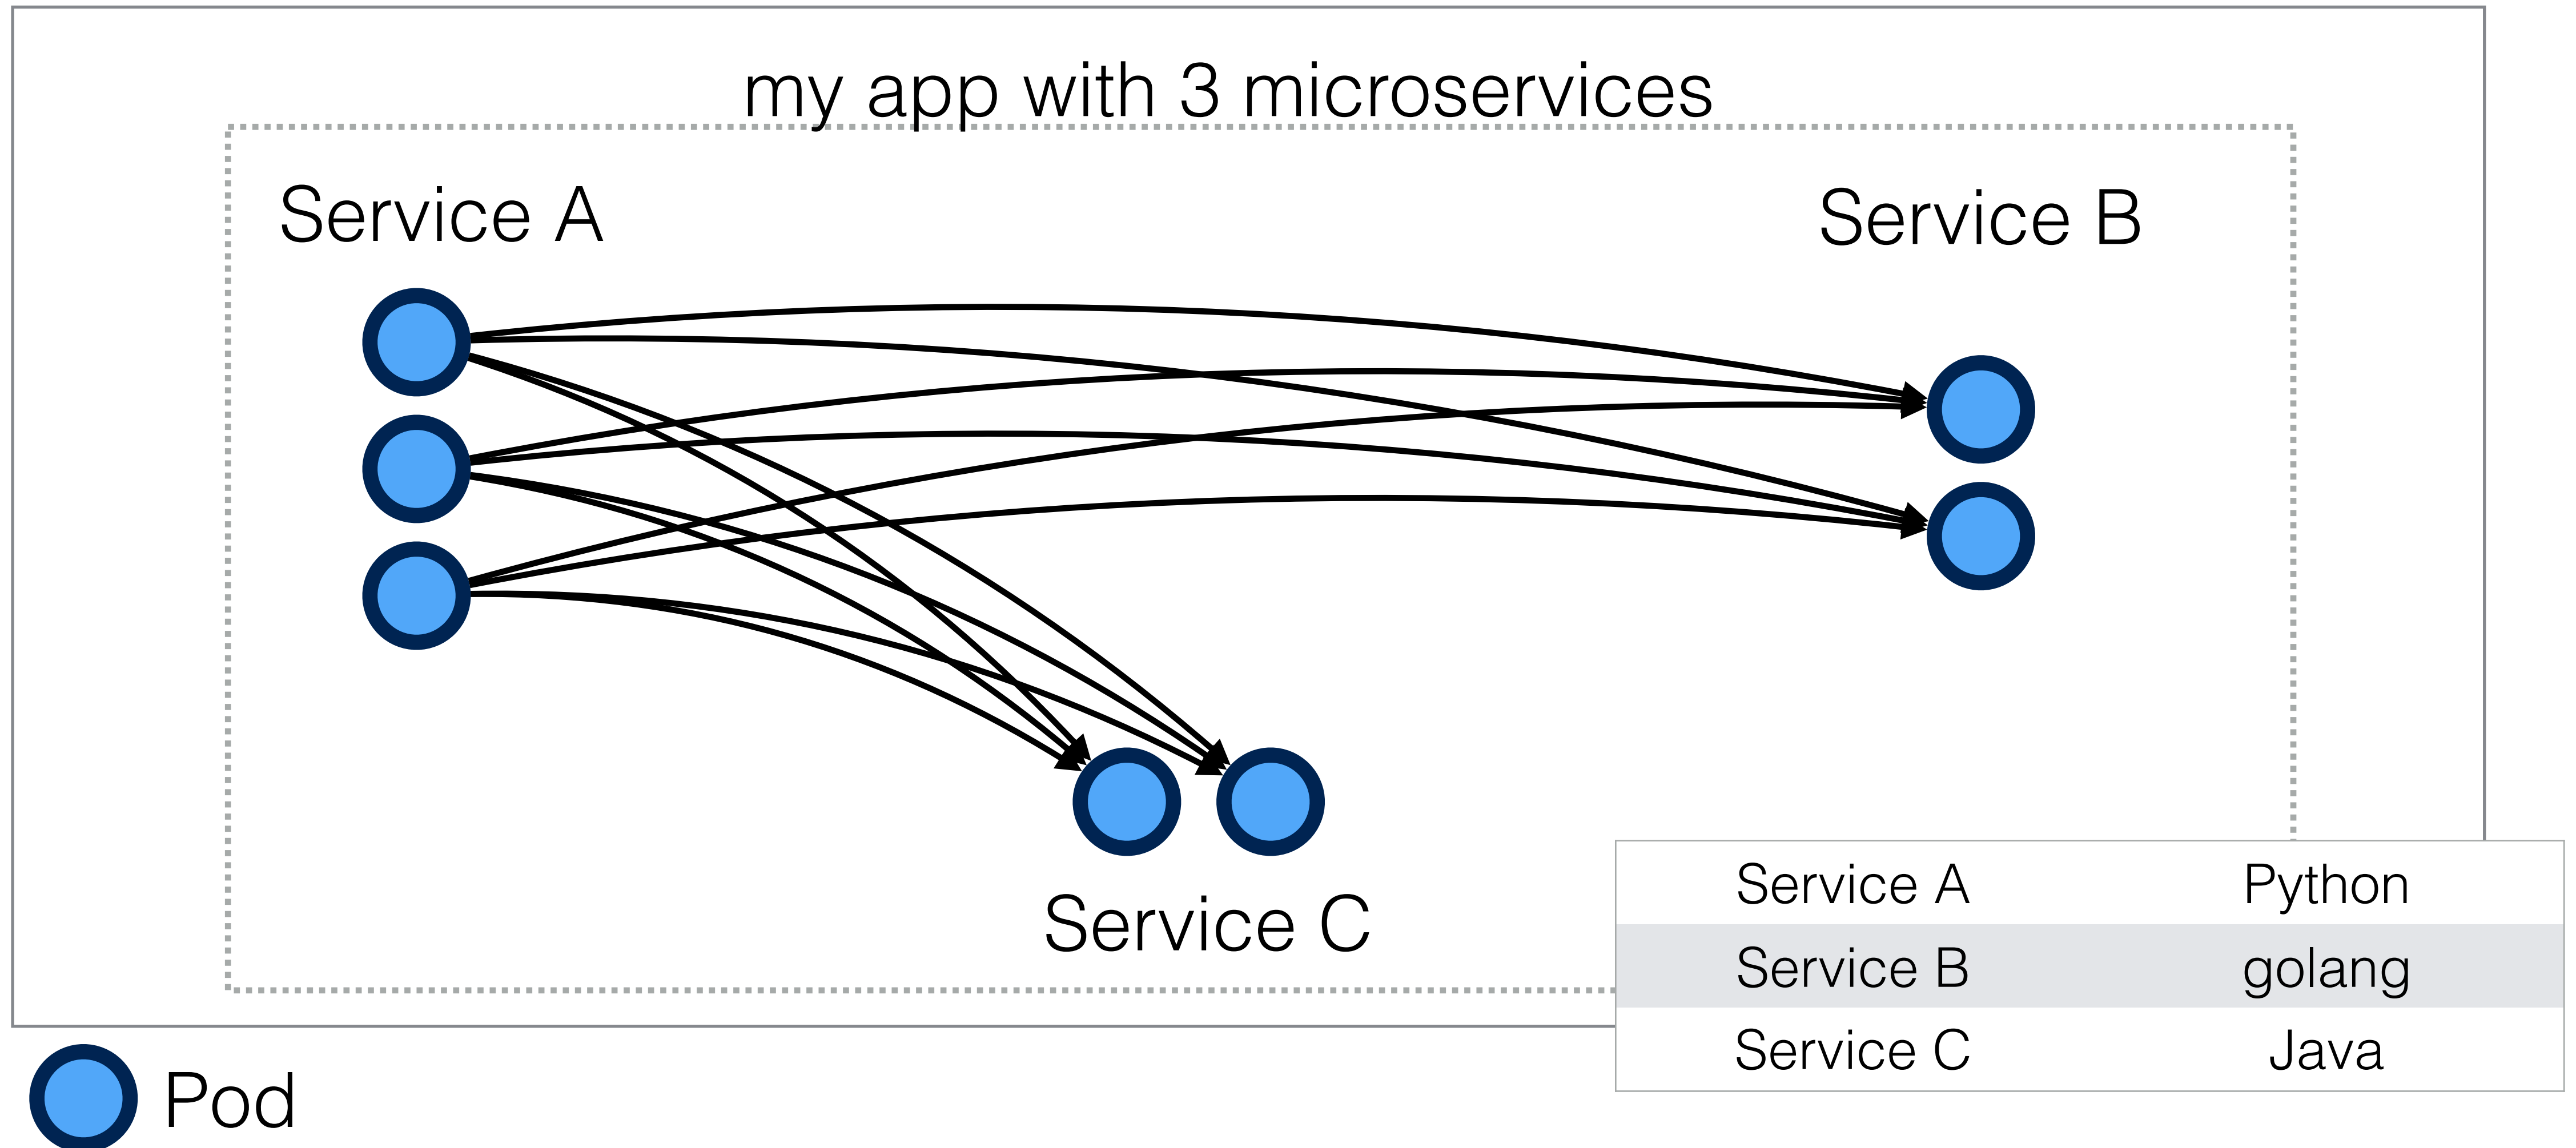
# Istio Components

---

- Pilot (control plane)
  - Responsible for **service discovery**, **traffic management** and **resiliency**
    - A/B tests and canary deployments
    - Timeouts, retries, circuit breakers
  - It does this by converting Istio rules to Envoy configurations
- Istio Auth (control plane)
  - Service-to-service and end-user **authentication** using **mutual TLS**

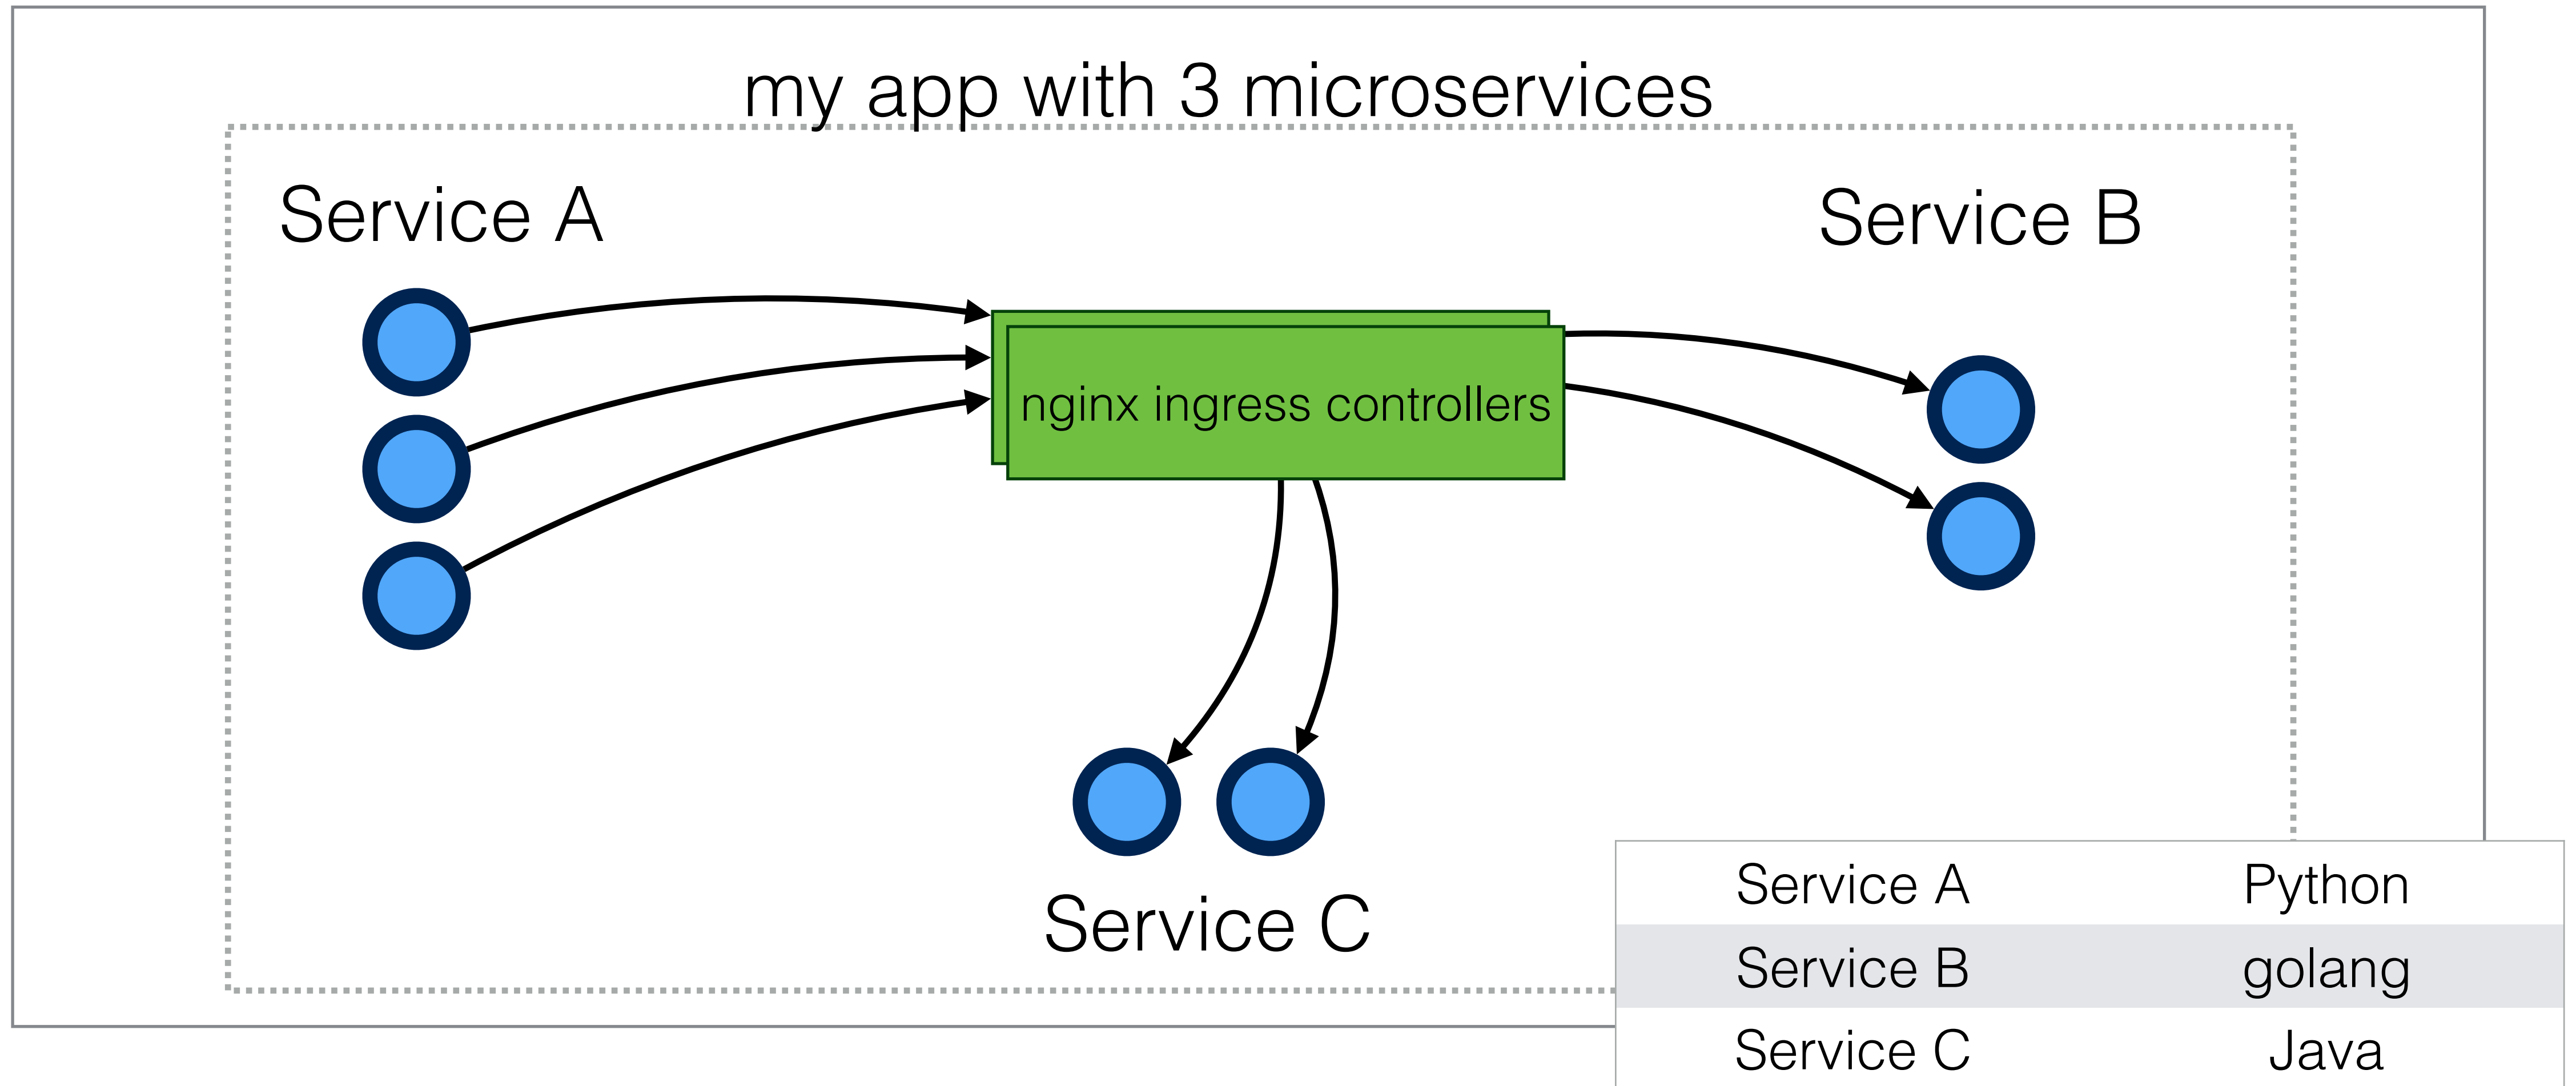
# Service Mesh

Kubernetes cluster



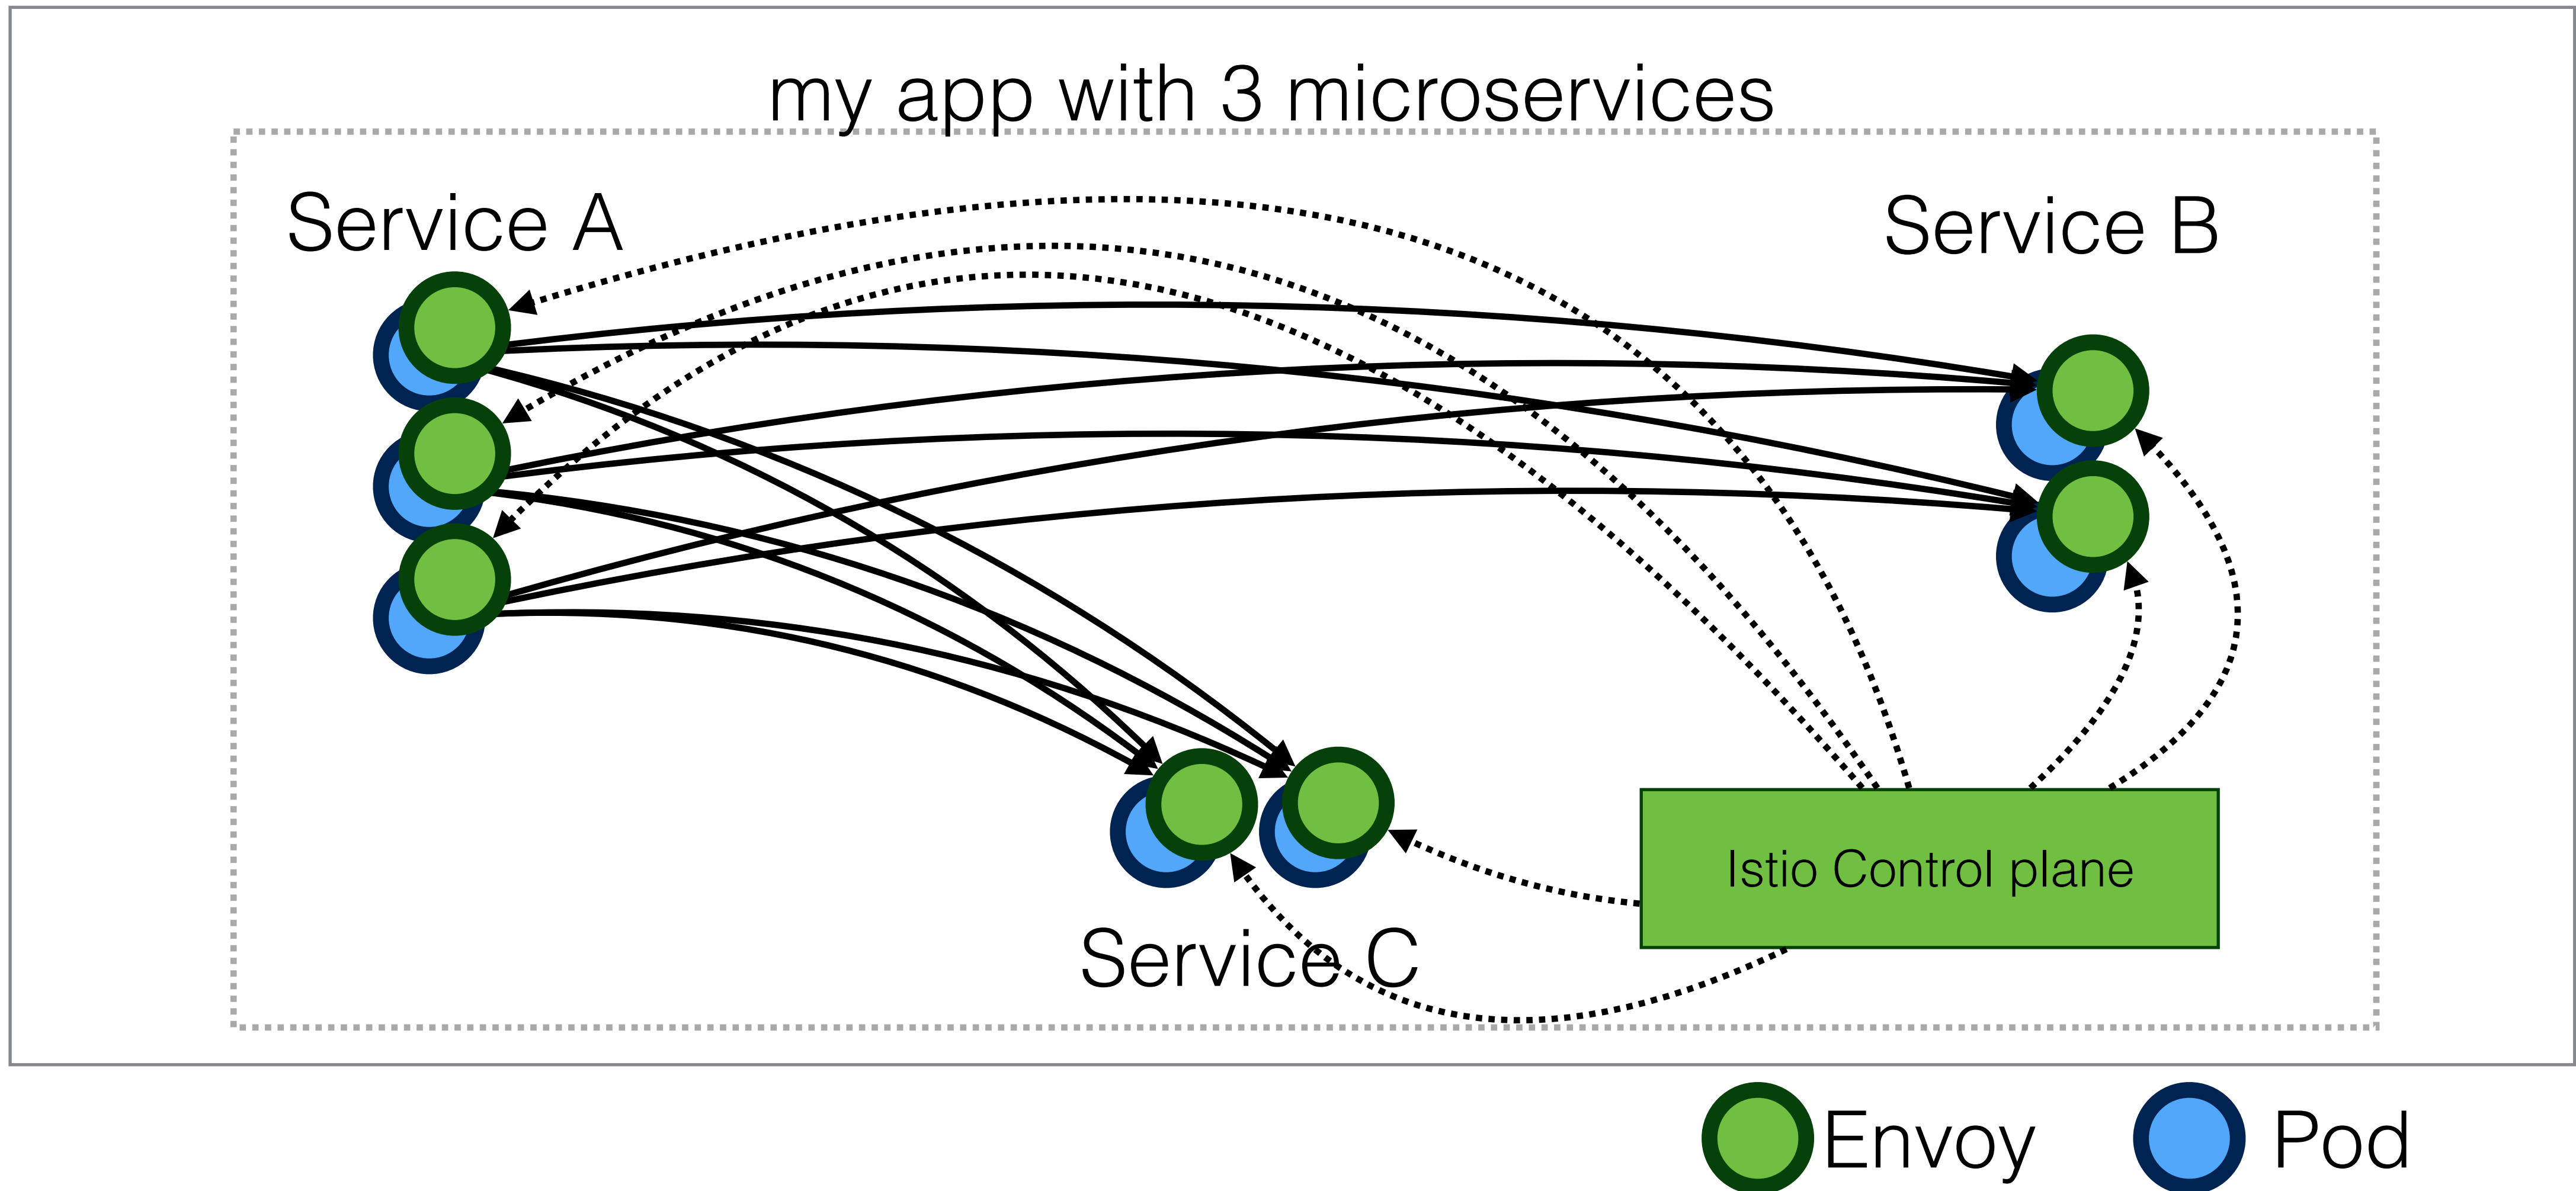
# Service Mesh - ingress

Kubernetes cluster



# Service Mesh - istio

Kubernetes cluster





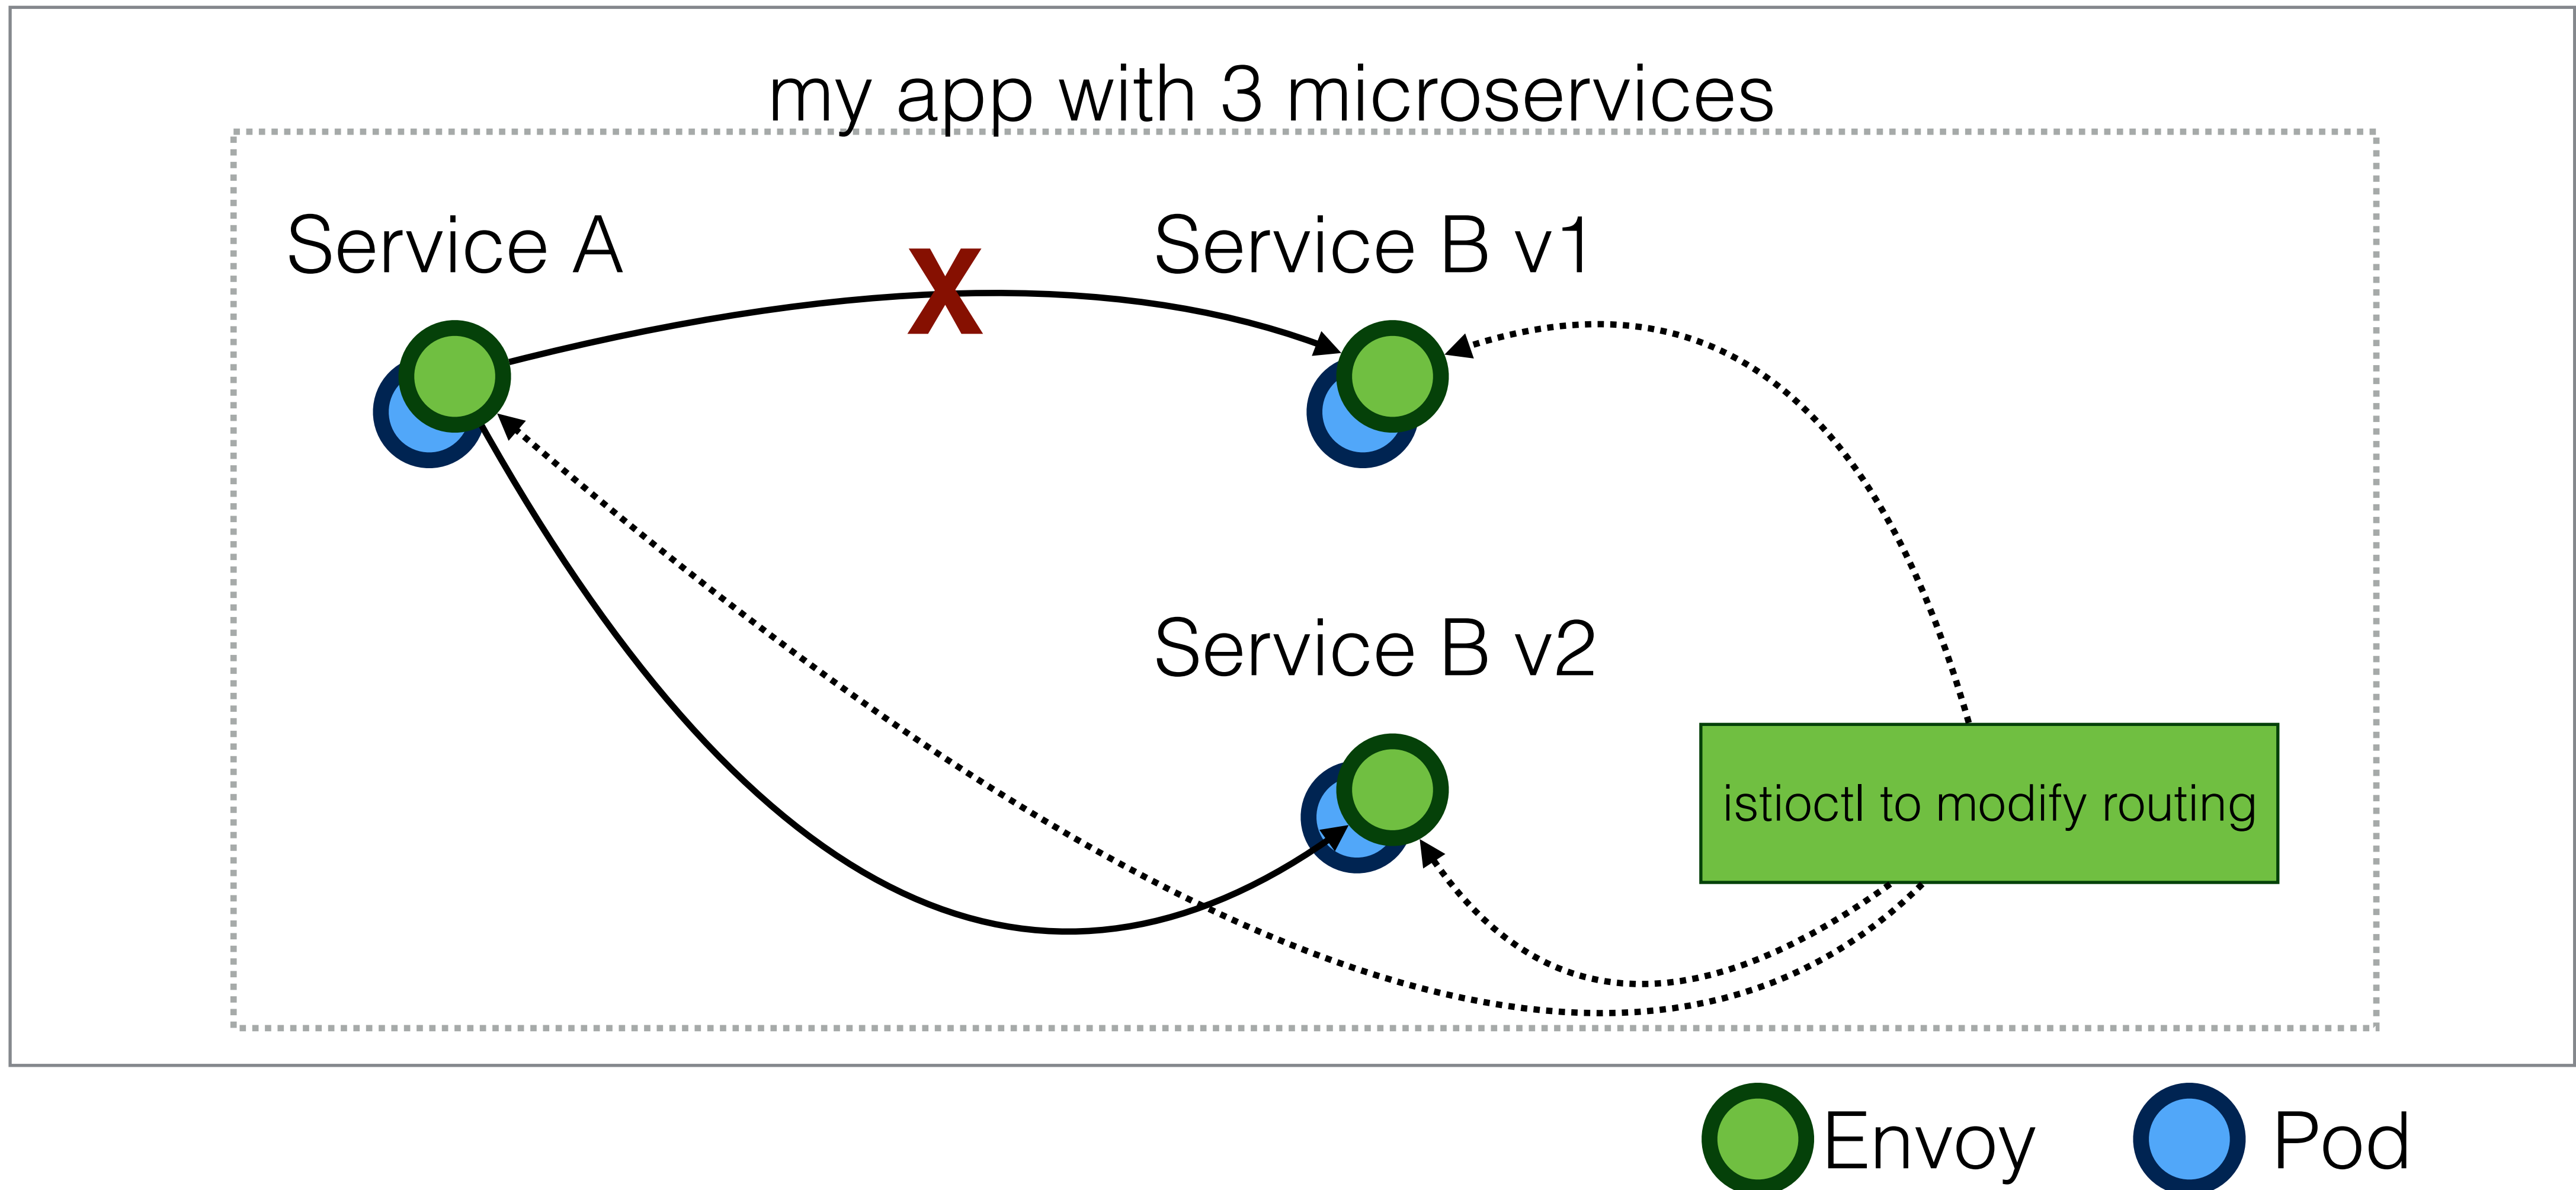
# Demo

istio demo



# traffic management

Kubernetes cluster



# Demo

istio demo - traffic management


# Demo

istio demo - distributed tracing

Calico

# Calico

- Calico provides **secure network connectivity** for **containers** and virtual machine workloads. (Definition: <https://docs.projectcalico.org/v3.1/introduction/>)



# Calico

---

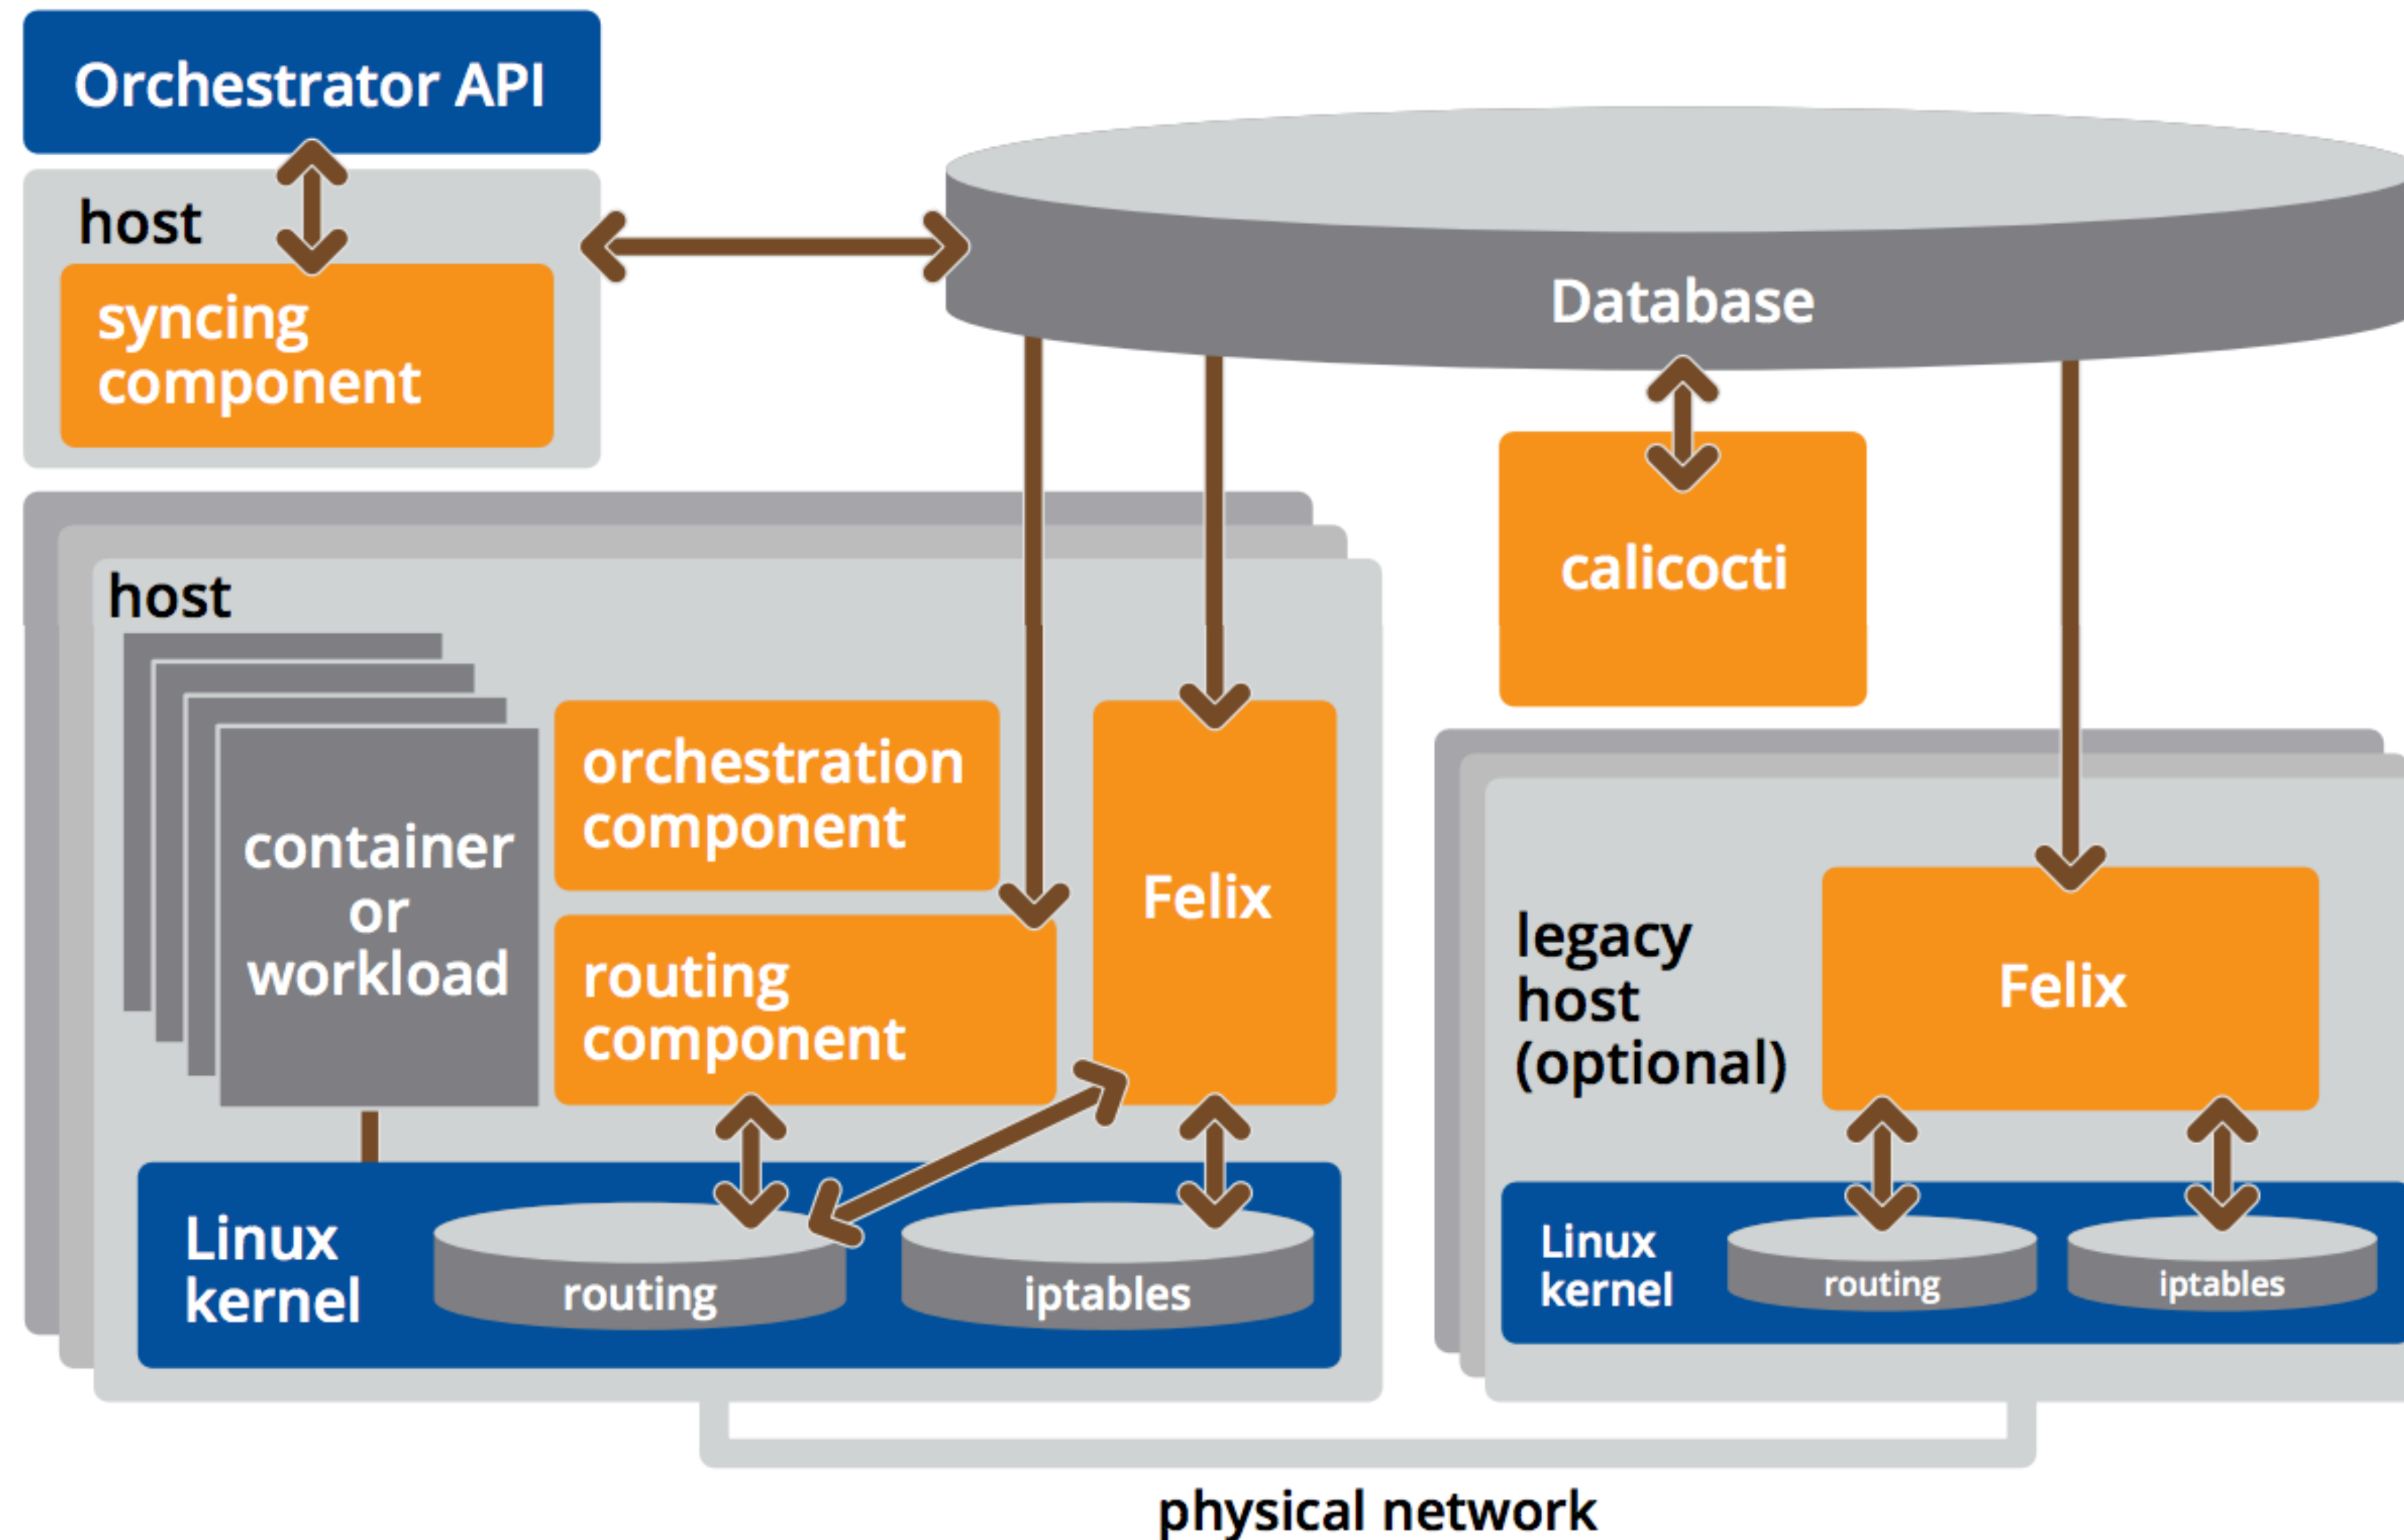
- Calico is a **Software Defined Network**, with a simplified model, with cloud-native in mind
- Calico creates a flat **Layer 3 network** using **BGP** (Border Gateway Protocol) as **routing mechanism**
  - BGP is also used as the “internet routing protocol” to route between providers (it’s a proven, scalable technology)
- **Policy driven network security** using the **Kubernetes Network Policy API**
  - Fine-grain control over the network, using the same Kubernetes API (using yaml files) as you’re used to
- **Only use overlay if necessary**, reducing overhead and increasing performance
  - An overlay network does **IP encapsulation**, but often those IP packets can be **routed without adding those extra headers** to IP packets

# Calico

---

- Works with **Kubernetes**, but also with **OpenStack**, **Mesos**, and others
- Uses **etcd as backend** (Kubernetes also uses etcd - a distributed key value store using Raft consensus)
- Works on **major cloud providers** like AWS, GCE (also Kubernetes Engine), Azure (ACS)
  - Will also support the hosted kubernetes services AWS EKS and Azure AKS when they'll be GA
- Works well within **enterprise environments**
  - Either **without overlay**
  - With **IP-in-IP tunneling**
  - Or using an **overlay** (VxLAN) network like **Flannel**

# Calico



Source: <https://www.projectcalico.org/wp-content/uploads/2018/01/ProjectCalico.v3.datasheet.pdf>



# Calico

---

- Calicoctl
  - Allows you to **manage the Calico network** and **security policy**
- Felix
  - **Daemon** that runs on every machine (calico-node **DaemonSet**)
  - Responsible for
    - **programming routes and ACL** on the nodes itself
    - **Interface management** (interacts with kernel - think about MAC address / IP level configuration)
  - Reports on **health** and **state** of the network

# Calico

---

- **BGP Client** (BIRD)
  - Runs **next to Felix** (still within the calico-node DaemonSet)
  - **Reads routing state** that Felix programmed and **distributes this information to other nodes**
    - Basically that's what BGP needs to do, it needs to make the **other nodes aware of routing information** to ensure **traffic** is **efficiently routed**
- **BGP Route Reflector**
  - All BGP clients are connected to each other, which may become a limiting factor
  - In larger deployments, a BGP route reflector might be setup, which acts as a **central point where BGP clients connect to** (instead of having a mesh topology)

# Calico

---

- Once Calico is setup, you can **create a network policy in Kubernetes**
- You can first create a **network policy to deny all access to all pods** (then afterwards you can open the ports that are needed):

```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: deny-all
  namespace: apps
spec:
  podSelector:
    matchLabels: {}
```

- At this point the pods are isolated, you'll not be able to connect from one pod to another anymore

# Calico

---

- **Isolated vs non-isolated**

- By default pods are **non-isolated**
  - Pods **accept traffic from any source**
- By having a **network policy** with a **selector** that selects them (the previous one selects all pods), network access is denied by default
  - The pod now becomes **isolated**
  - Only connections that are defined in the **network policy** are allowed
- This is on a namespace basis

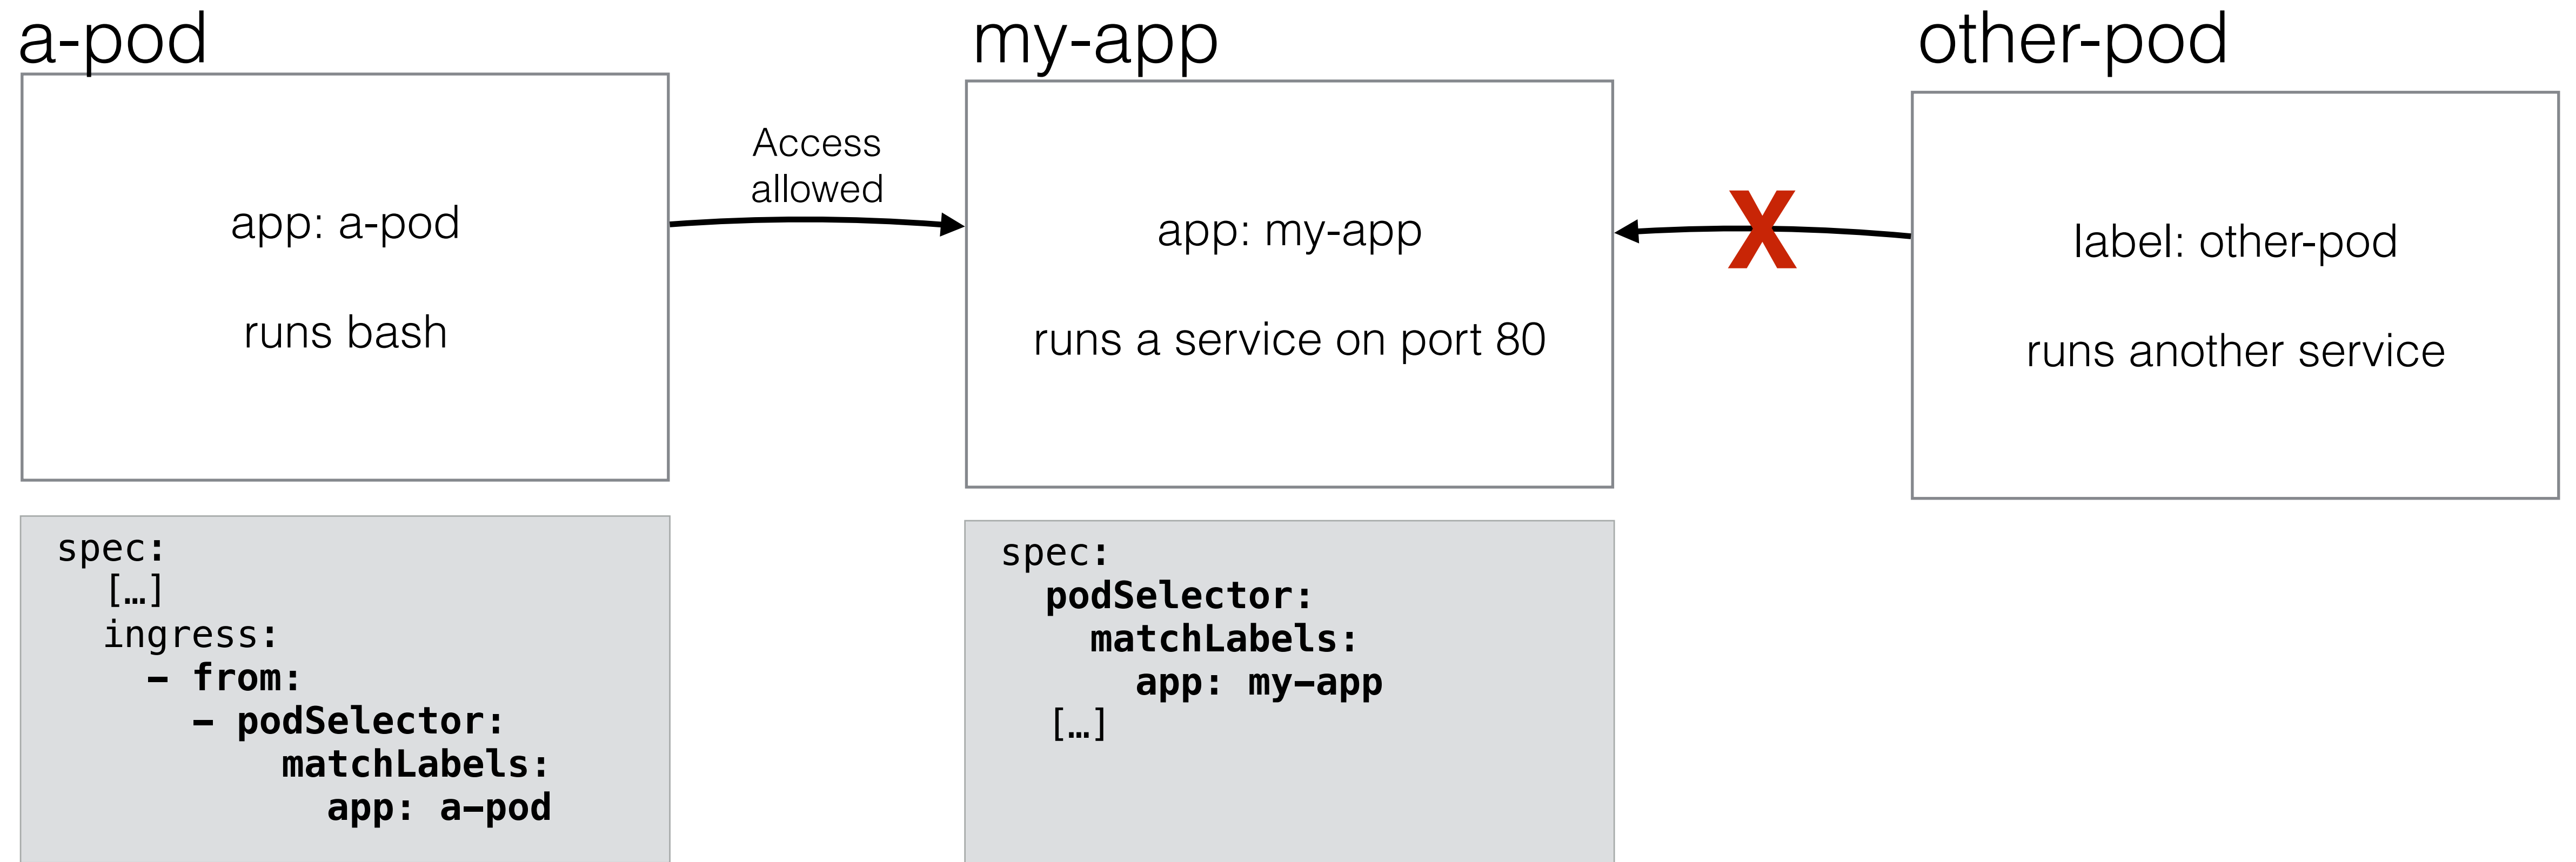
# Calico

---

- You can now add a new rule to enable network access to a pod:

```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: allow-my-app
  namespace: apps
spec:
  podSelector:
    matchLabels:
      app: my-app
  ingress:
    - from:
      - podSelector:
          matchLabels:
            app: a-pod
```

# Calico



# Demo

Calico example

# Demo

Calico egress example



# Vault

Managing credentials in a distributed environment

# Vault

---

- Vault is a **tool for managing secrets**
  - For example: passwords, API keys, SSH keys, certificates
- It's **opensource** and released by **HashiCorp** (like Vagrant, terraform, and other well known tools)
- Some use cases are:
  - General Secret Storage
  - Employee Credential Storage (Sharing credentials, but using audit log, with ability to roll over credentials)
  - API key generation for scripts (Dynamic Secrets)
  - Data Encryption / Decryption

# Vault Features

---

- **Secure Secret Storage**
  - **Encrypted key-value pairs** can be stored in Vault
- **Dynamic Secrets**
  - Vault can create **on-demand secrets** and **revoke them after a period of time** (when the client lease is up)
  - For example **AWS credentials** to access an S3 bucket
- **Data Encryption**
  - Vault can encrypt / decrypt data without storing it

# Vault Features

---

- **Leasing and Renewal**

- Secrets in Vault have a **lease** (a time to live)
- When the lease is up, the **secret** will be **revoked** (deleted)
- Clients can ask for a **renewal** (a new secret) using an API

- **Revocation**

- Easy **revocation features**
- For example, all secrets of a **particular user** can be removed

# Vault Operator

---

- In April 2018 CoreOS released the Vault Operator
- It allows you to easily deploy Vault on Kubernetes
- It allows you to configure and maintain Vault using the Kubernetes API (using yaml files and kubectl)
- It gives you a **good alternative** to **secret management tools** on **public cloud** (like the AWS Secrets Manager or AWS Parameter store)

# Demo

Vault

# Openshift Origin

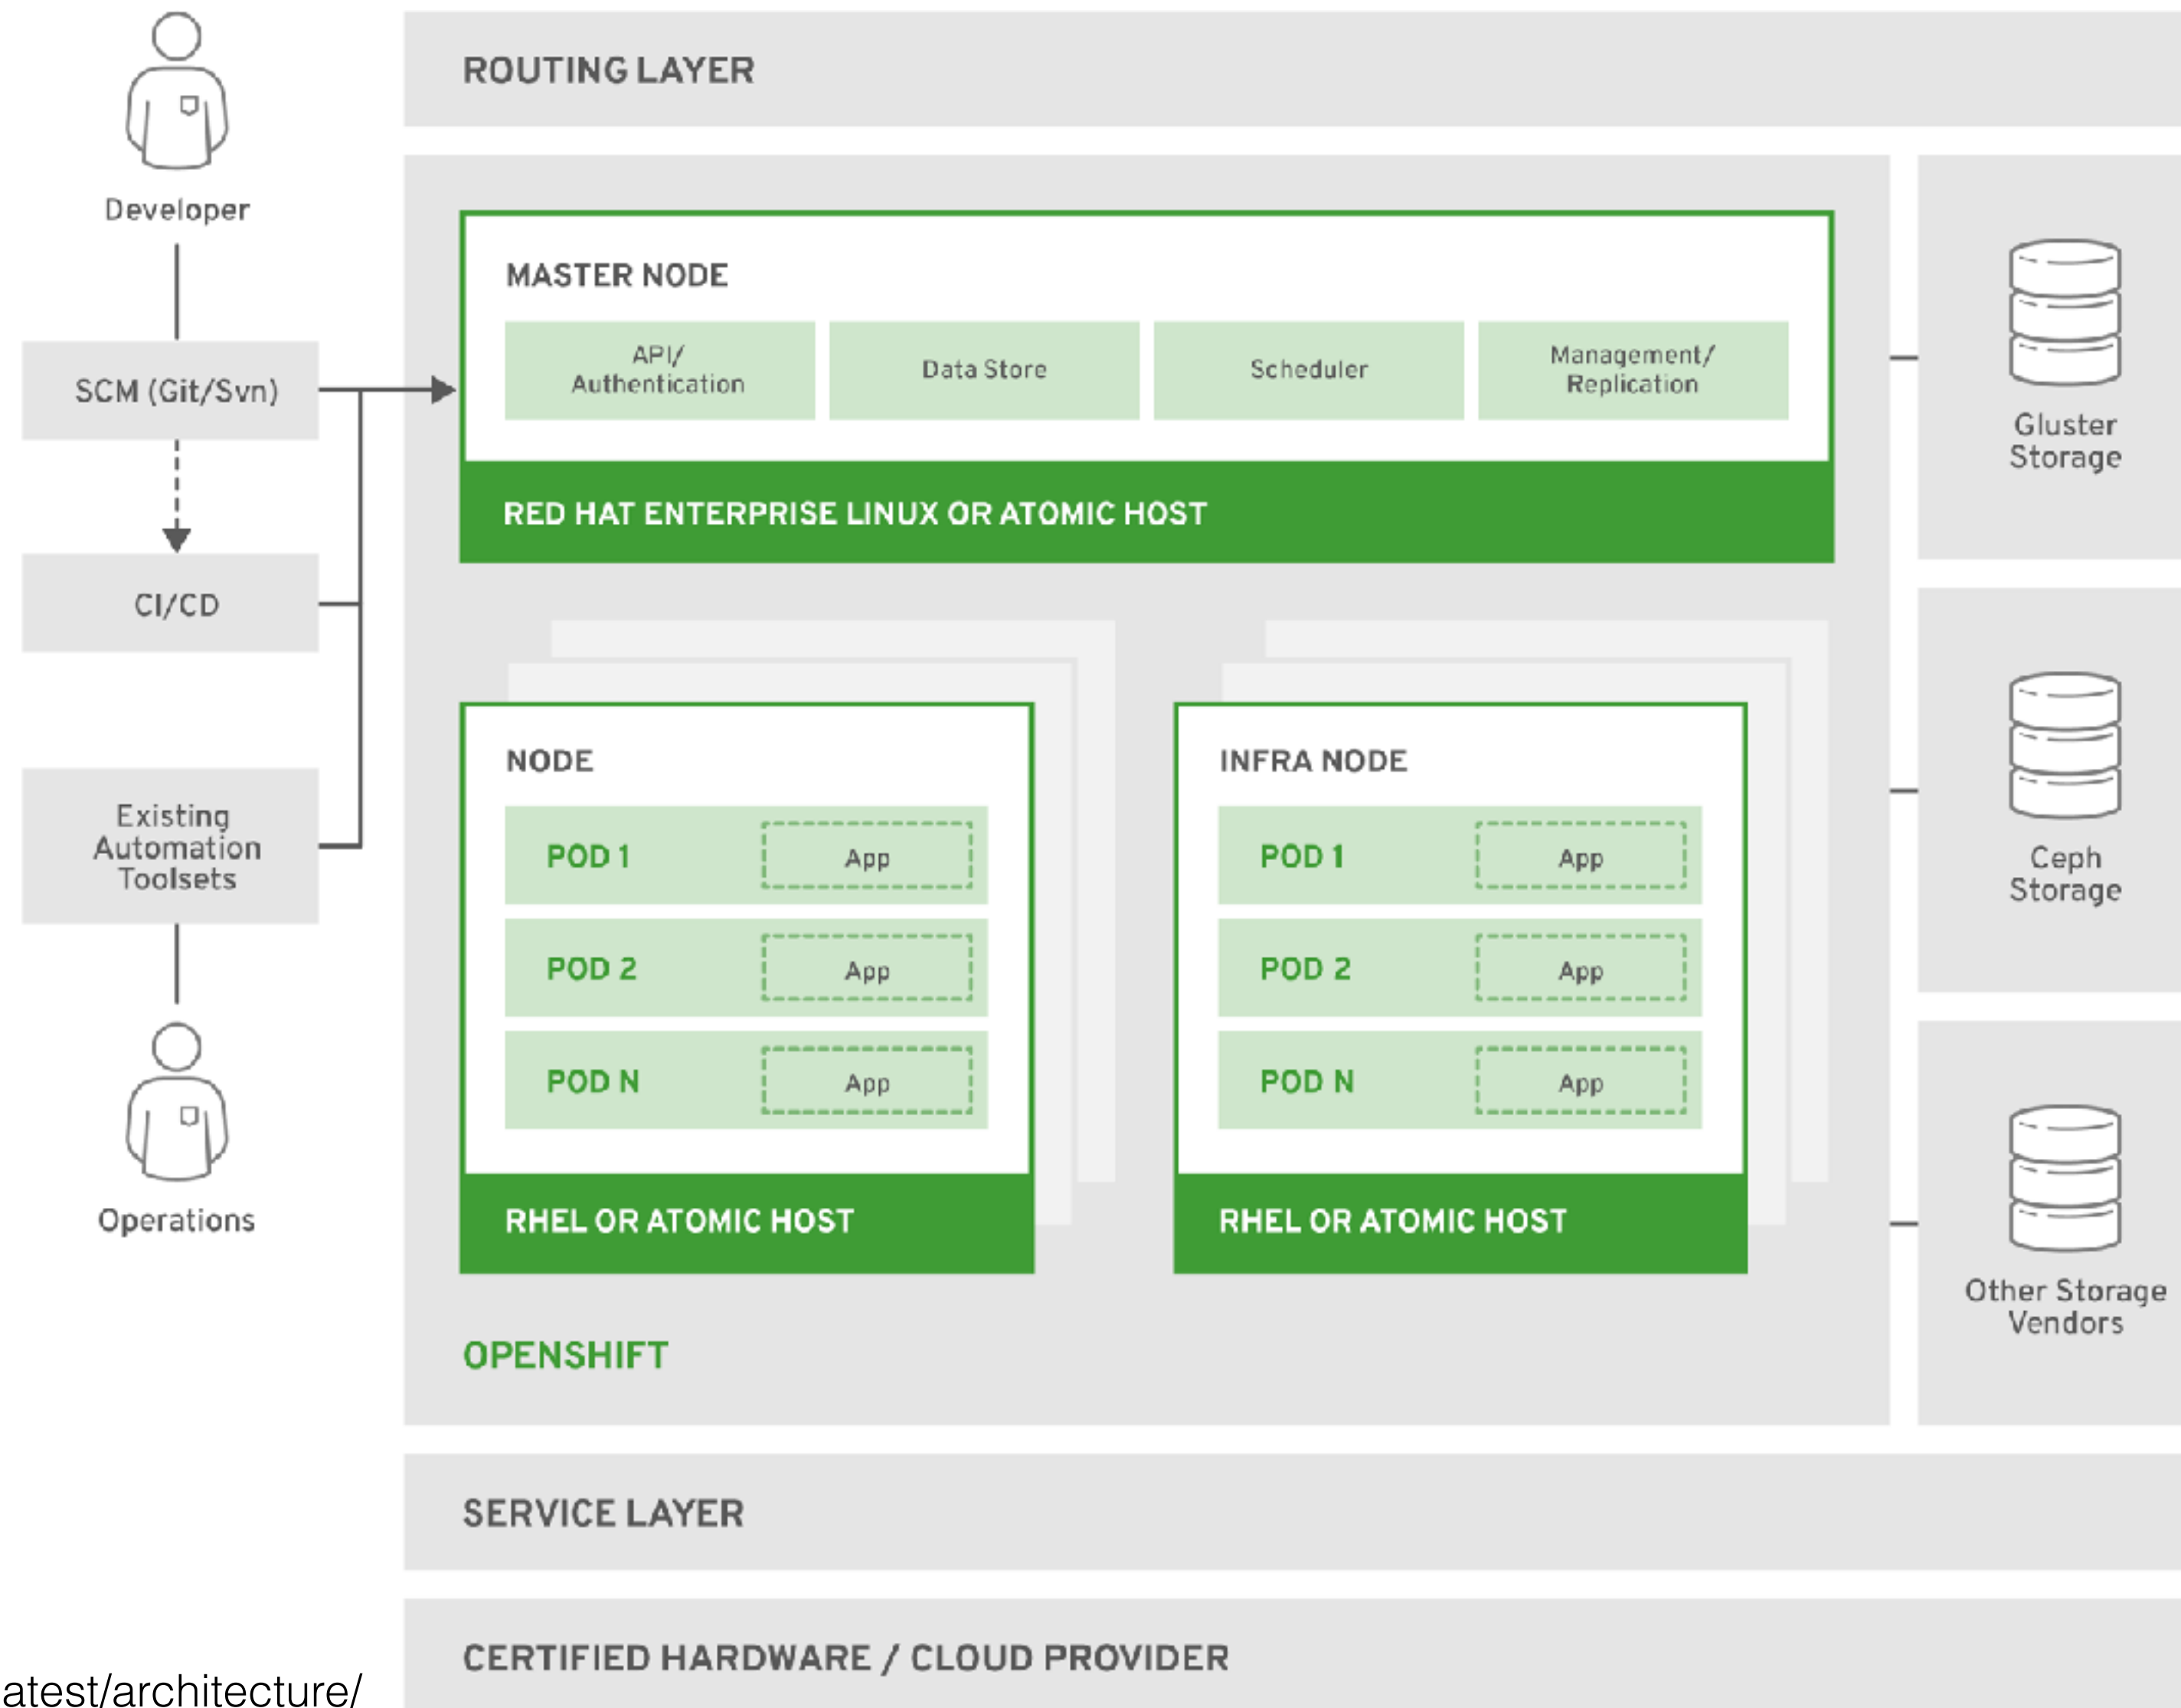
# OpenShift Origin

---

- OpenShift Origin is a **distribution of Kubernetes**
- It is optimized for **continuous application development** and **multi-tenancy**
- It adds **developer** and **operations** centric tools
- OpenShift Origin is the **upstream community project** that powers Openshift
- It uses **Kubernetes**, **Docker**, and **Project Atomic** (a container operating system)
- Definitions: <https://www.openshift.org/>



# OpenShift Origin Architecture



Source: <https://docs.openshift.org/latest/architecture/>

# OpenShift Origin

---

- Openshift also has a quick setup by running **openshift in a container**
- With “oc cluster up” you can bring up the Kubernetes cluster with the **Openshift frontend**
- Using the web frontend, you can **create projects and applications**
  - For example, you can start a **NodeJS** project or a **MySQL** database
  - Those projects will use a **git repository** to build and push the docker image
  - Everything happens **behind the scenes**, it's very developer focussed
- You can also **integrate with Jenkins** by putting Jenkinsfiles in your project

# OpenShift Origin

---

- The **Developer experience is great**, because it **hides** the **complexities** of Kubernetes
- That means that Openshift has its own **implementation** of:
  - Handling **storage** (for example ceph)
  - Handling **authentication** (you plugin into Openshift)
  - Integrating **CI, Ingress, loadbalancing, etc**
- This can be a pro or a con (you don't have to worry about the implementation, but you also have to follow Openshift's way of doing things)

# OpenShift Origin

---

- When to use Openshift?
  - If you need a **complete system** that **integrates CI/CD** and **Kubernetes** (and hosted platforms are no option)
  - If you don't want to design and develop a custom delivery platform for your developers, but are OK with using **Openshift's way of doing things**
  - If you “just want to **let your developers run apps on Kubernetes**”
  - If you **already are using Redhat**, and you'd like to get a on-prem PaaS offering with the support Redhat provides

# Demo

Installing OpenShift

# Demo

Running an application